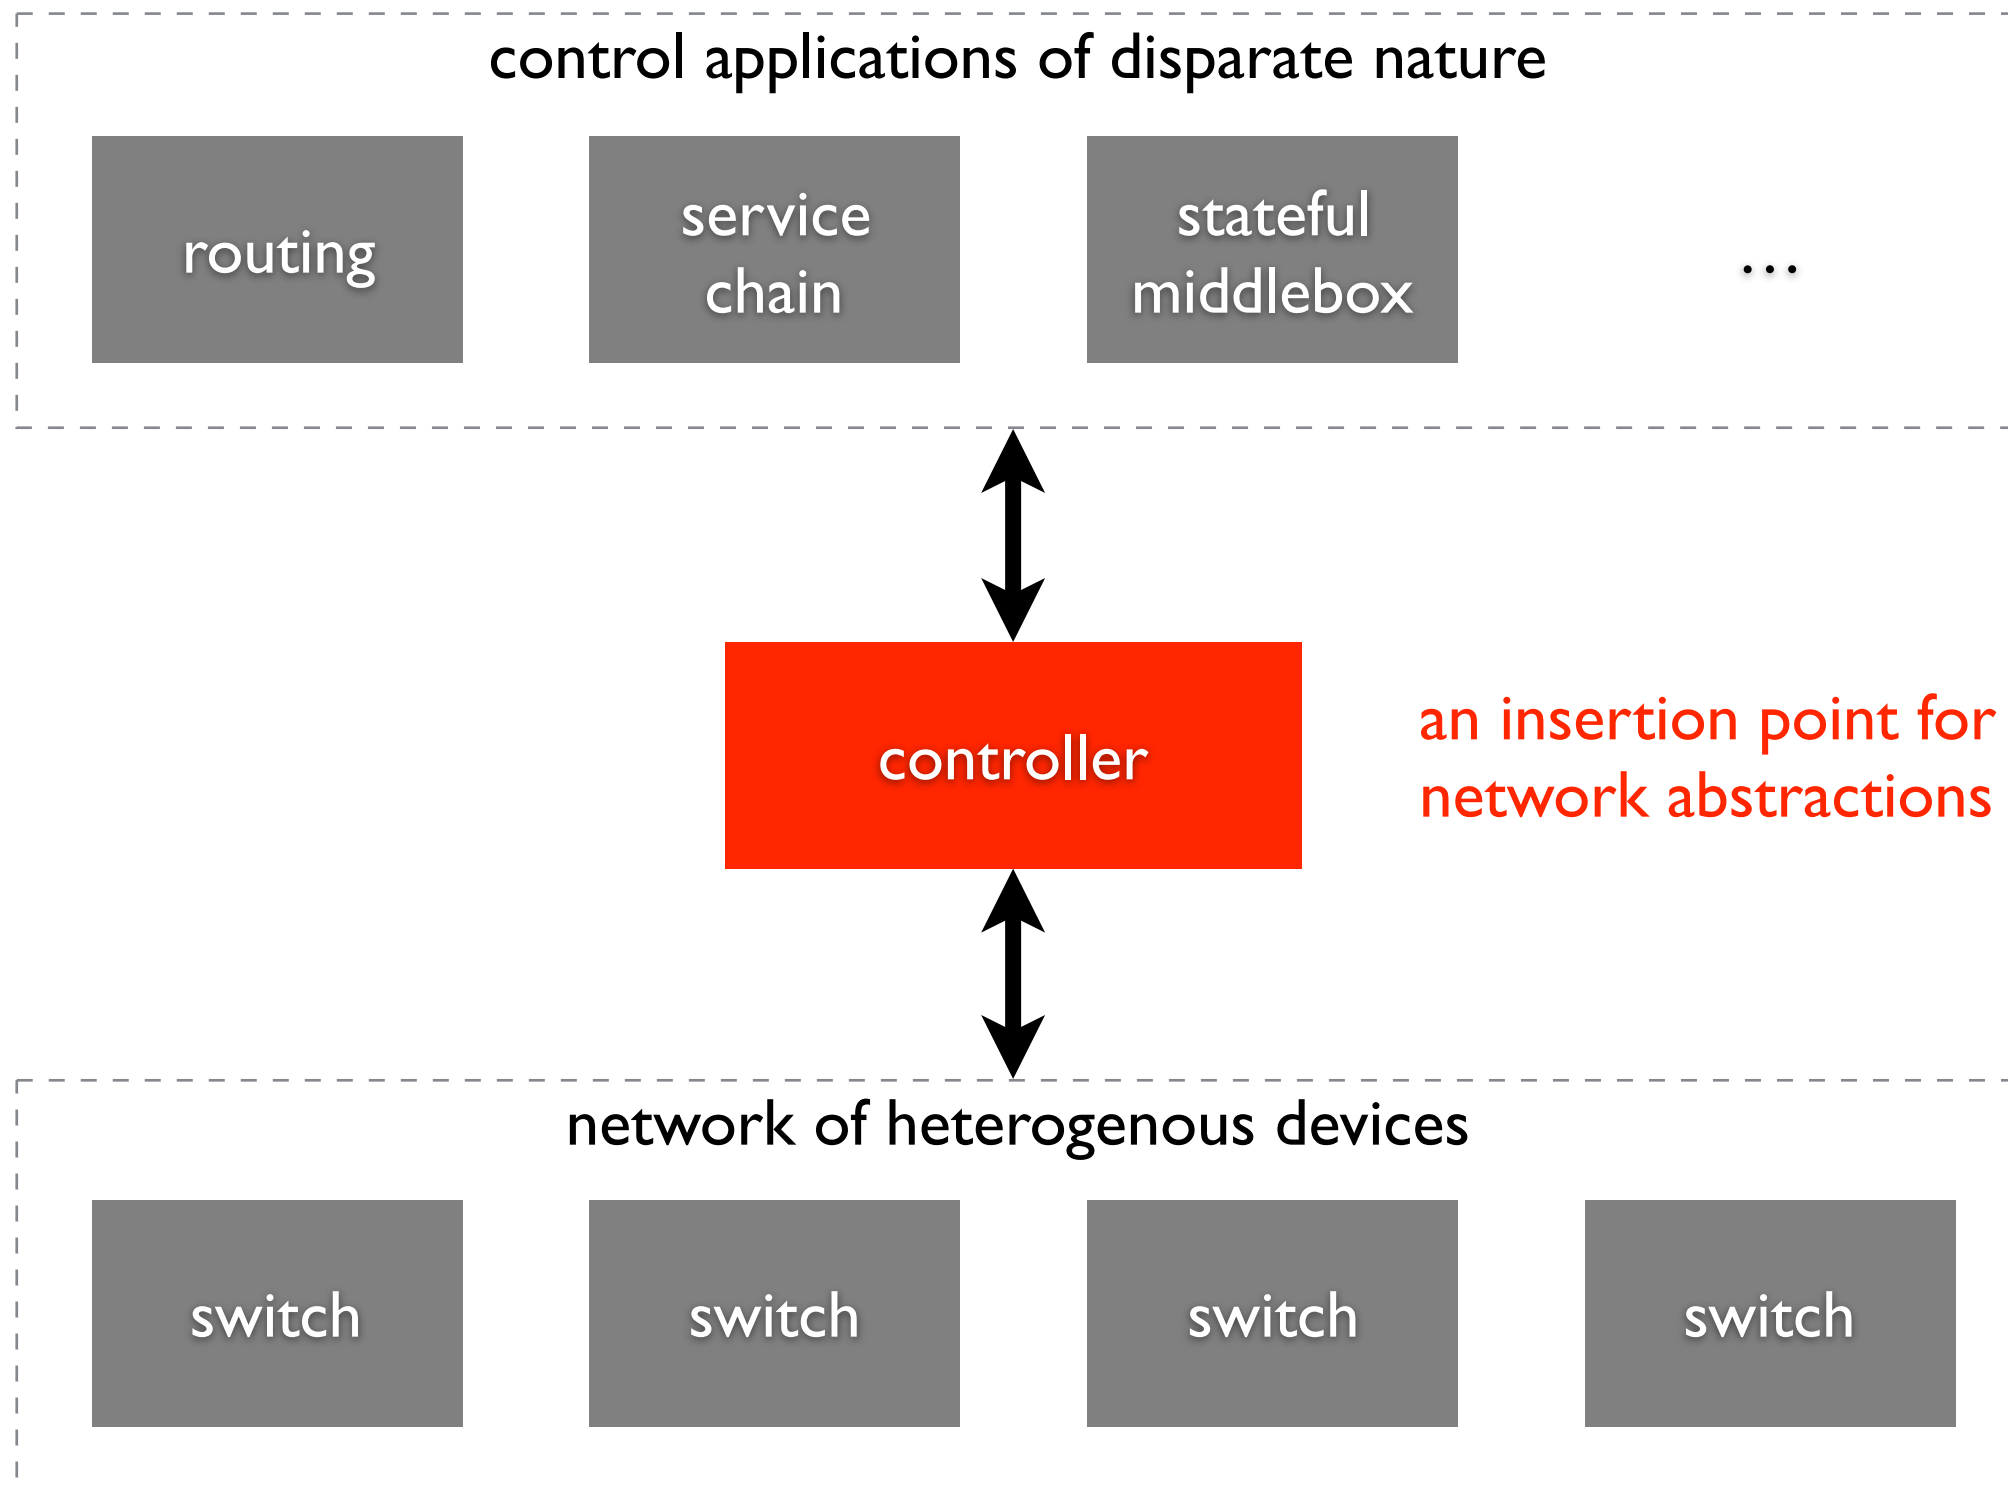




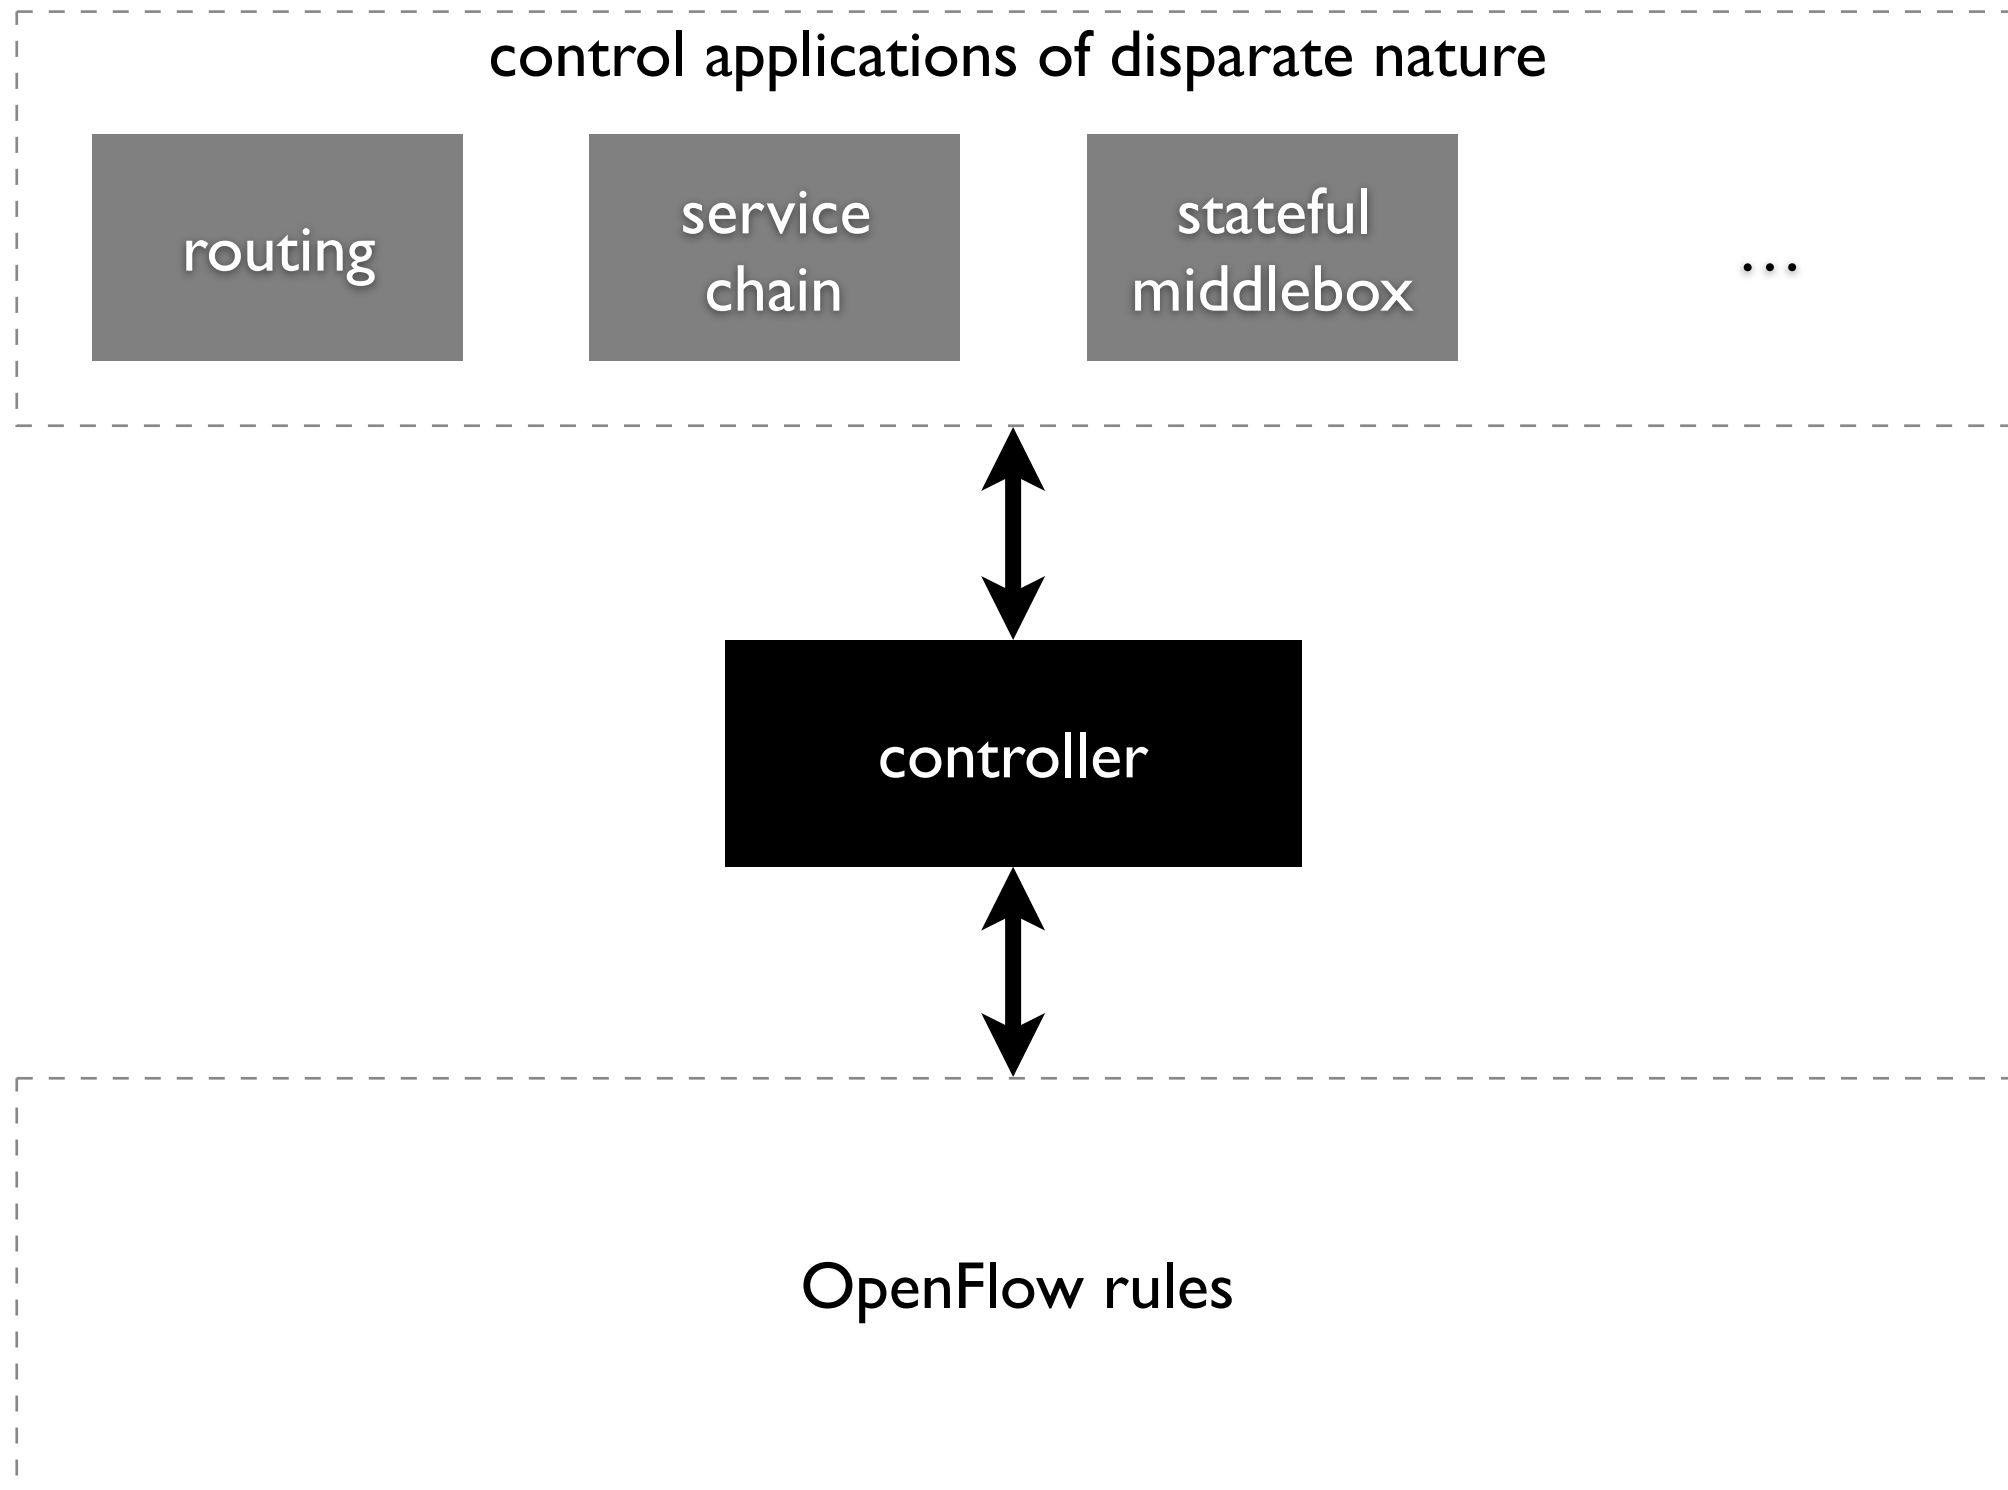
# *Ravel*: a database-defined network

Anduo Wang    Xueyuan Mei    Jason Croft  
Matthew Caesar    Brighten Godfrey

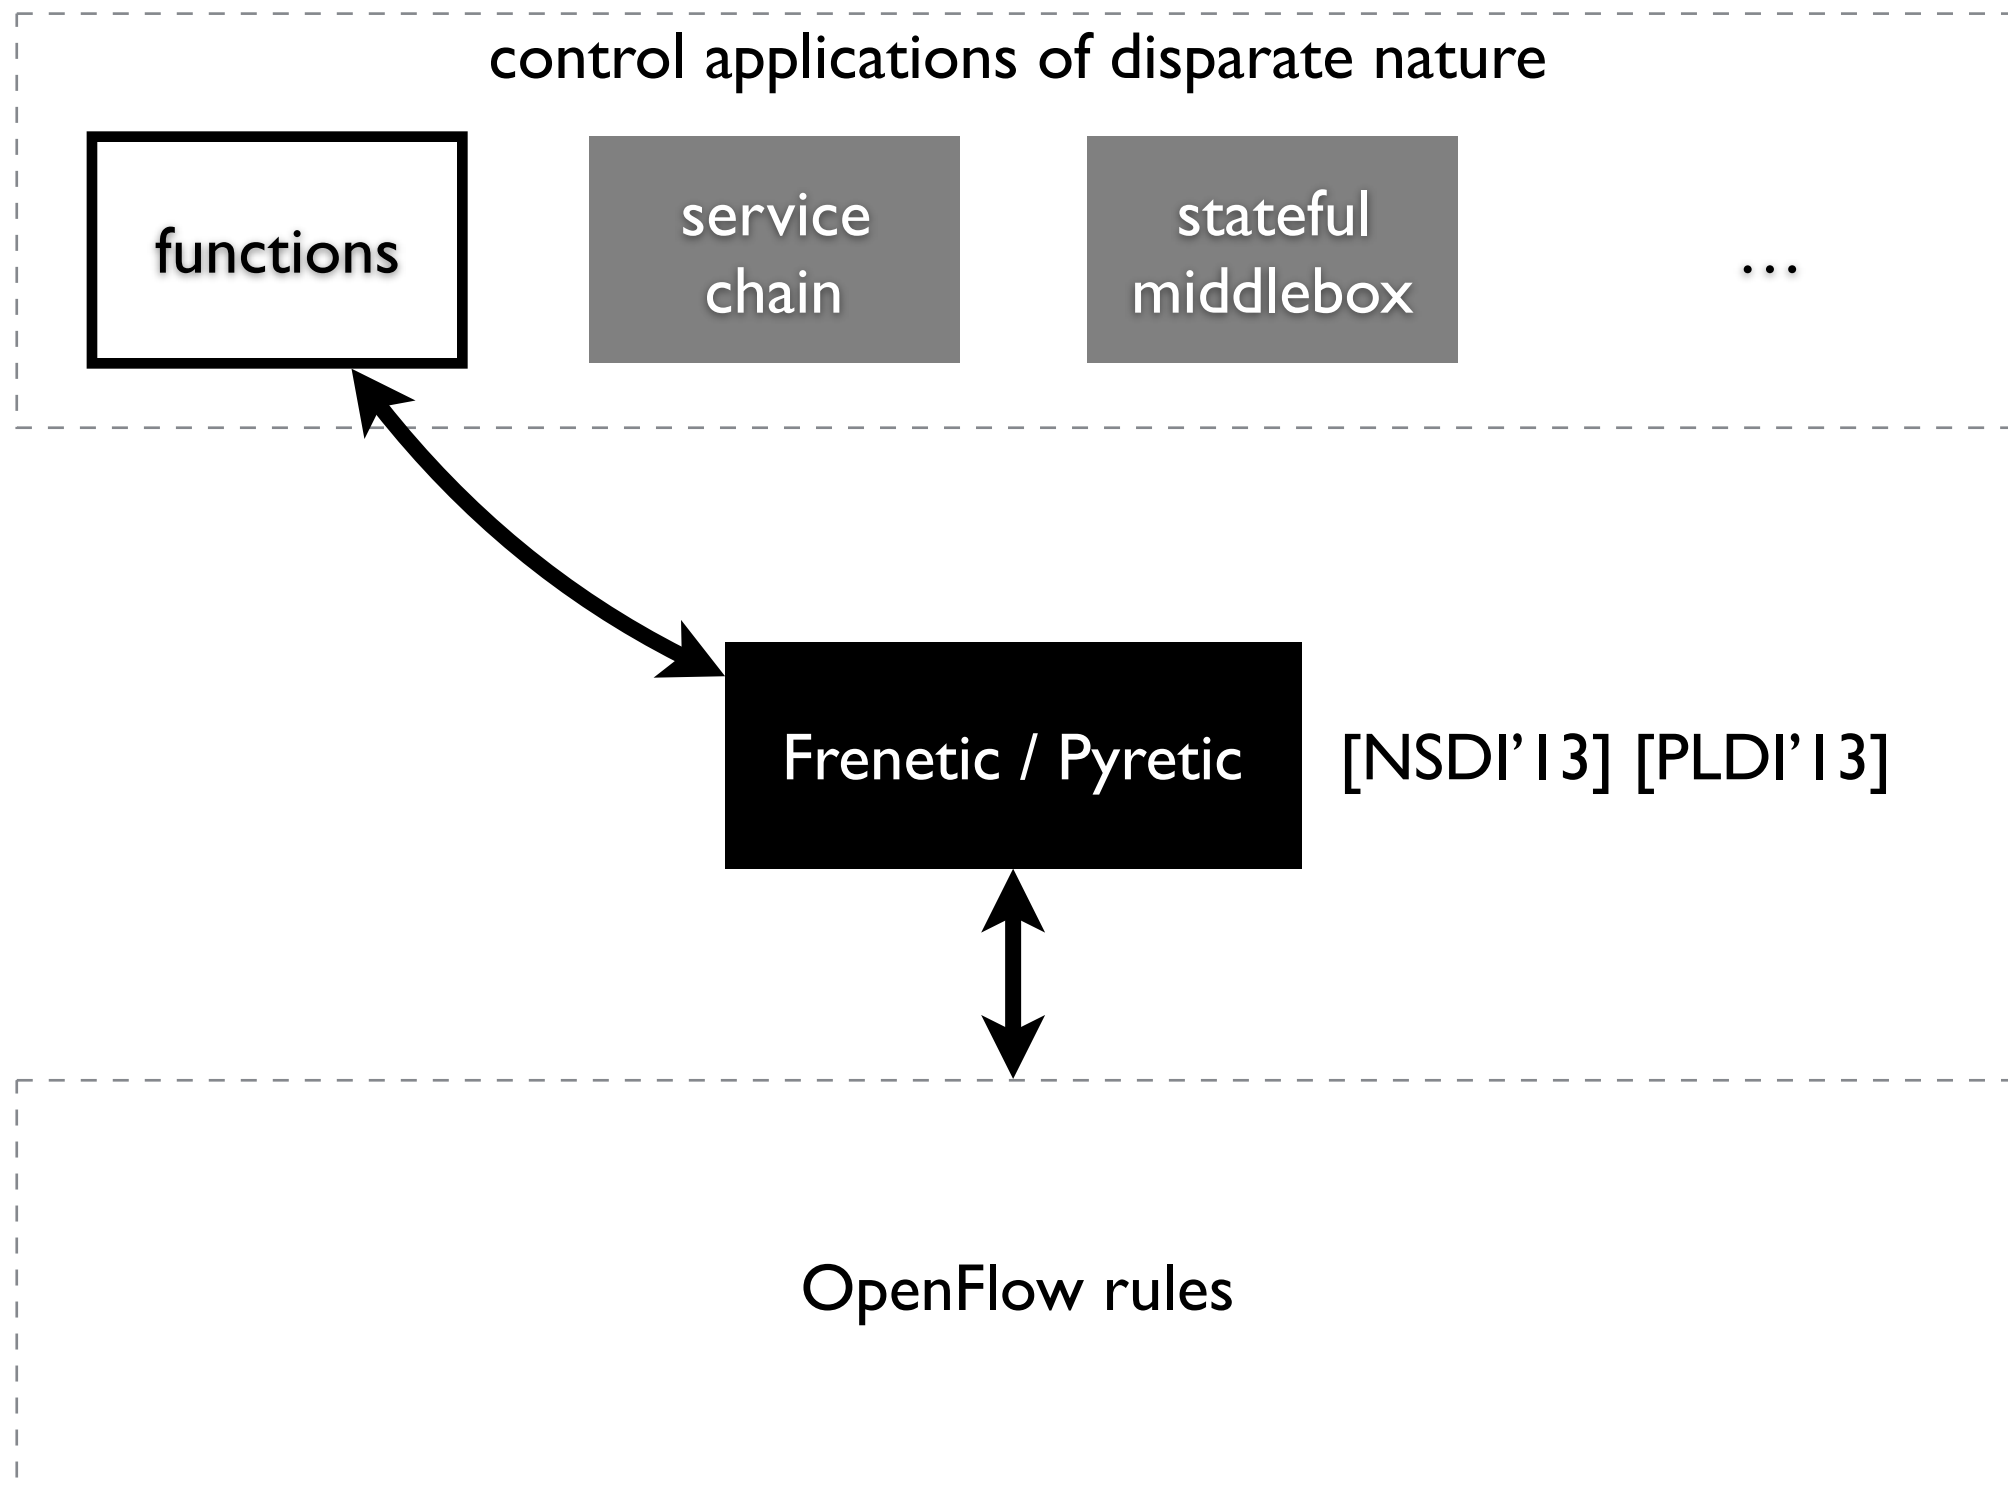
# software-defined network



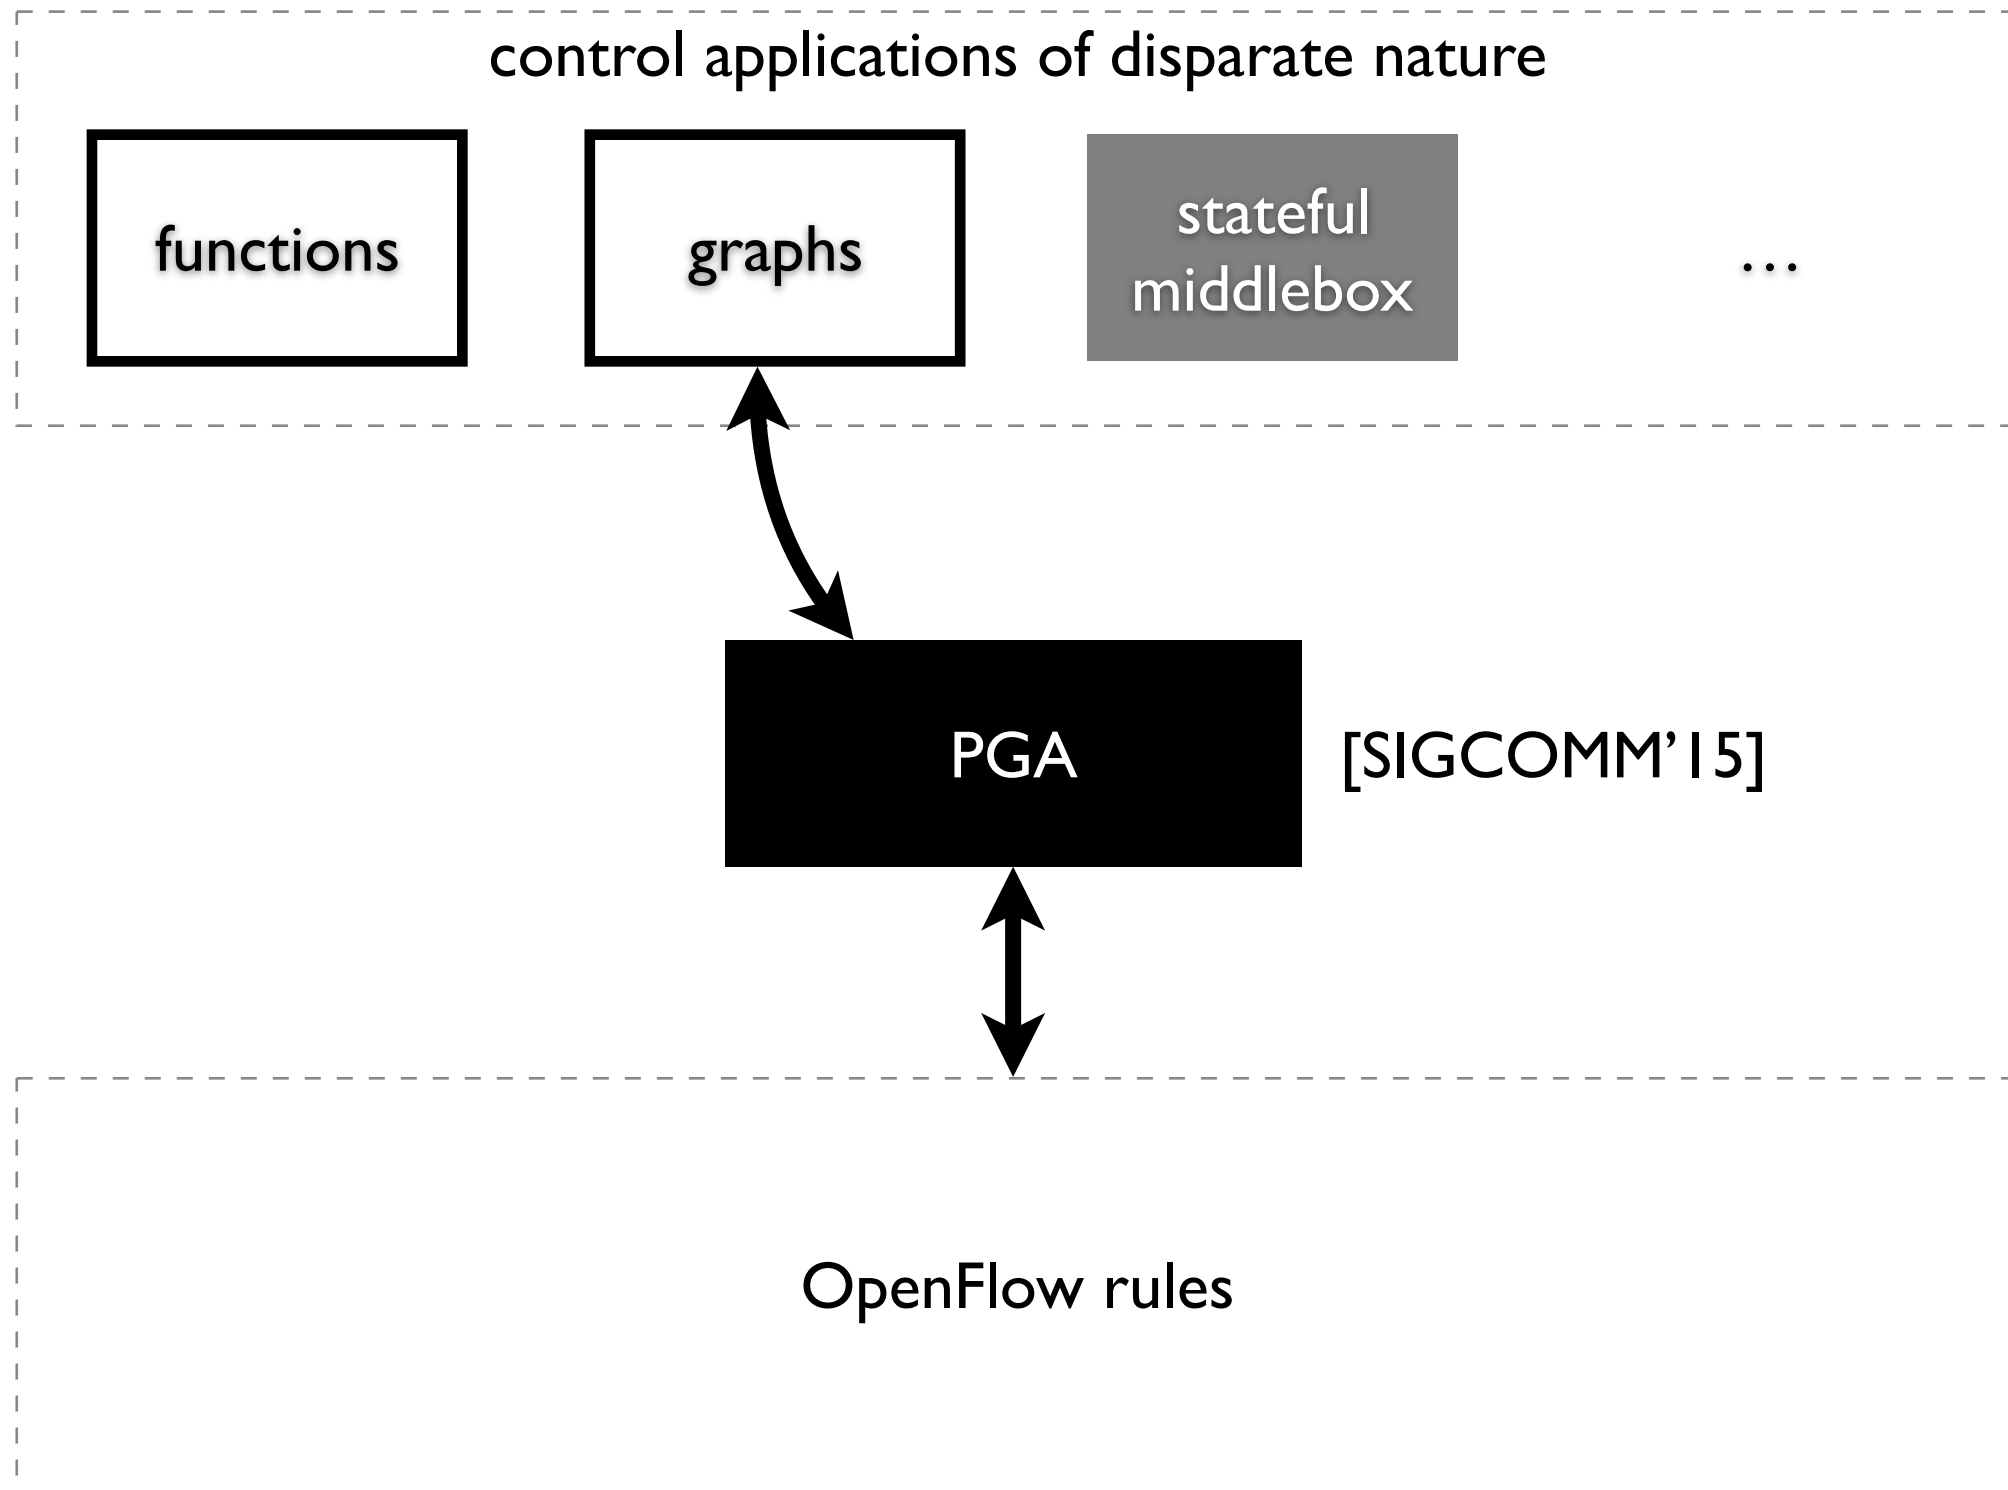
# abstractions



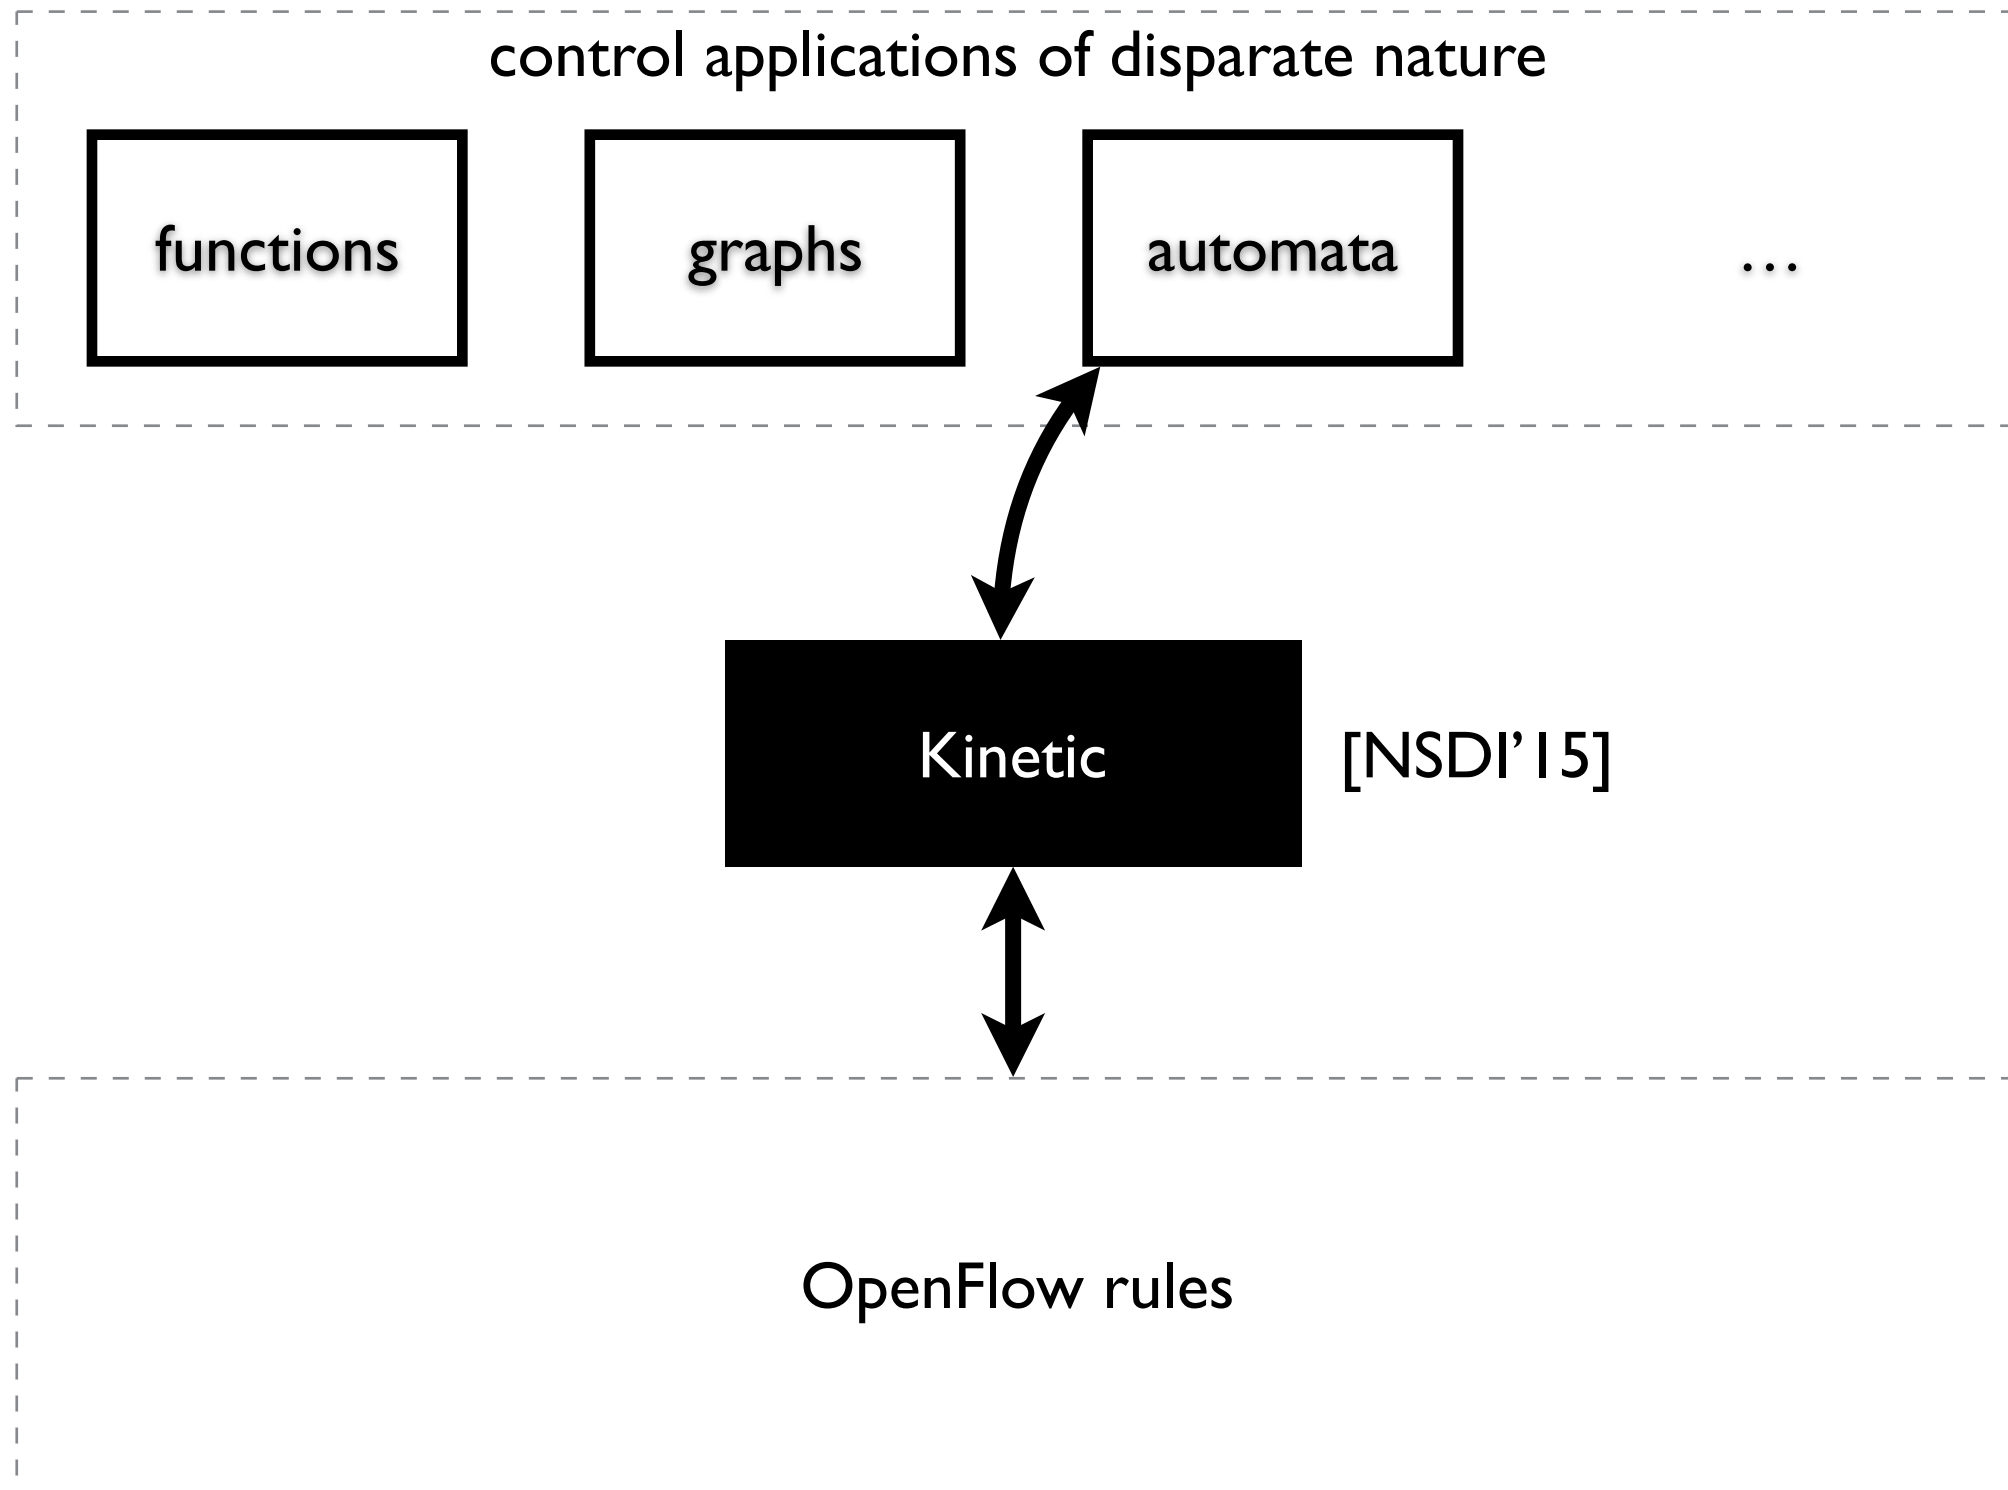
# abstractions



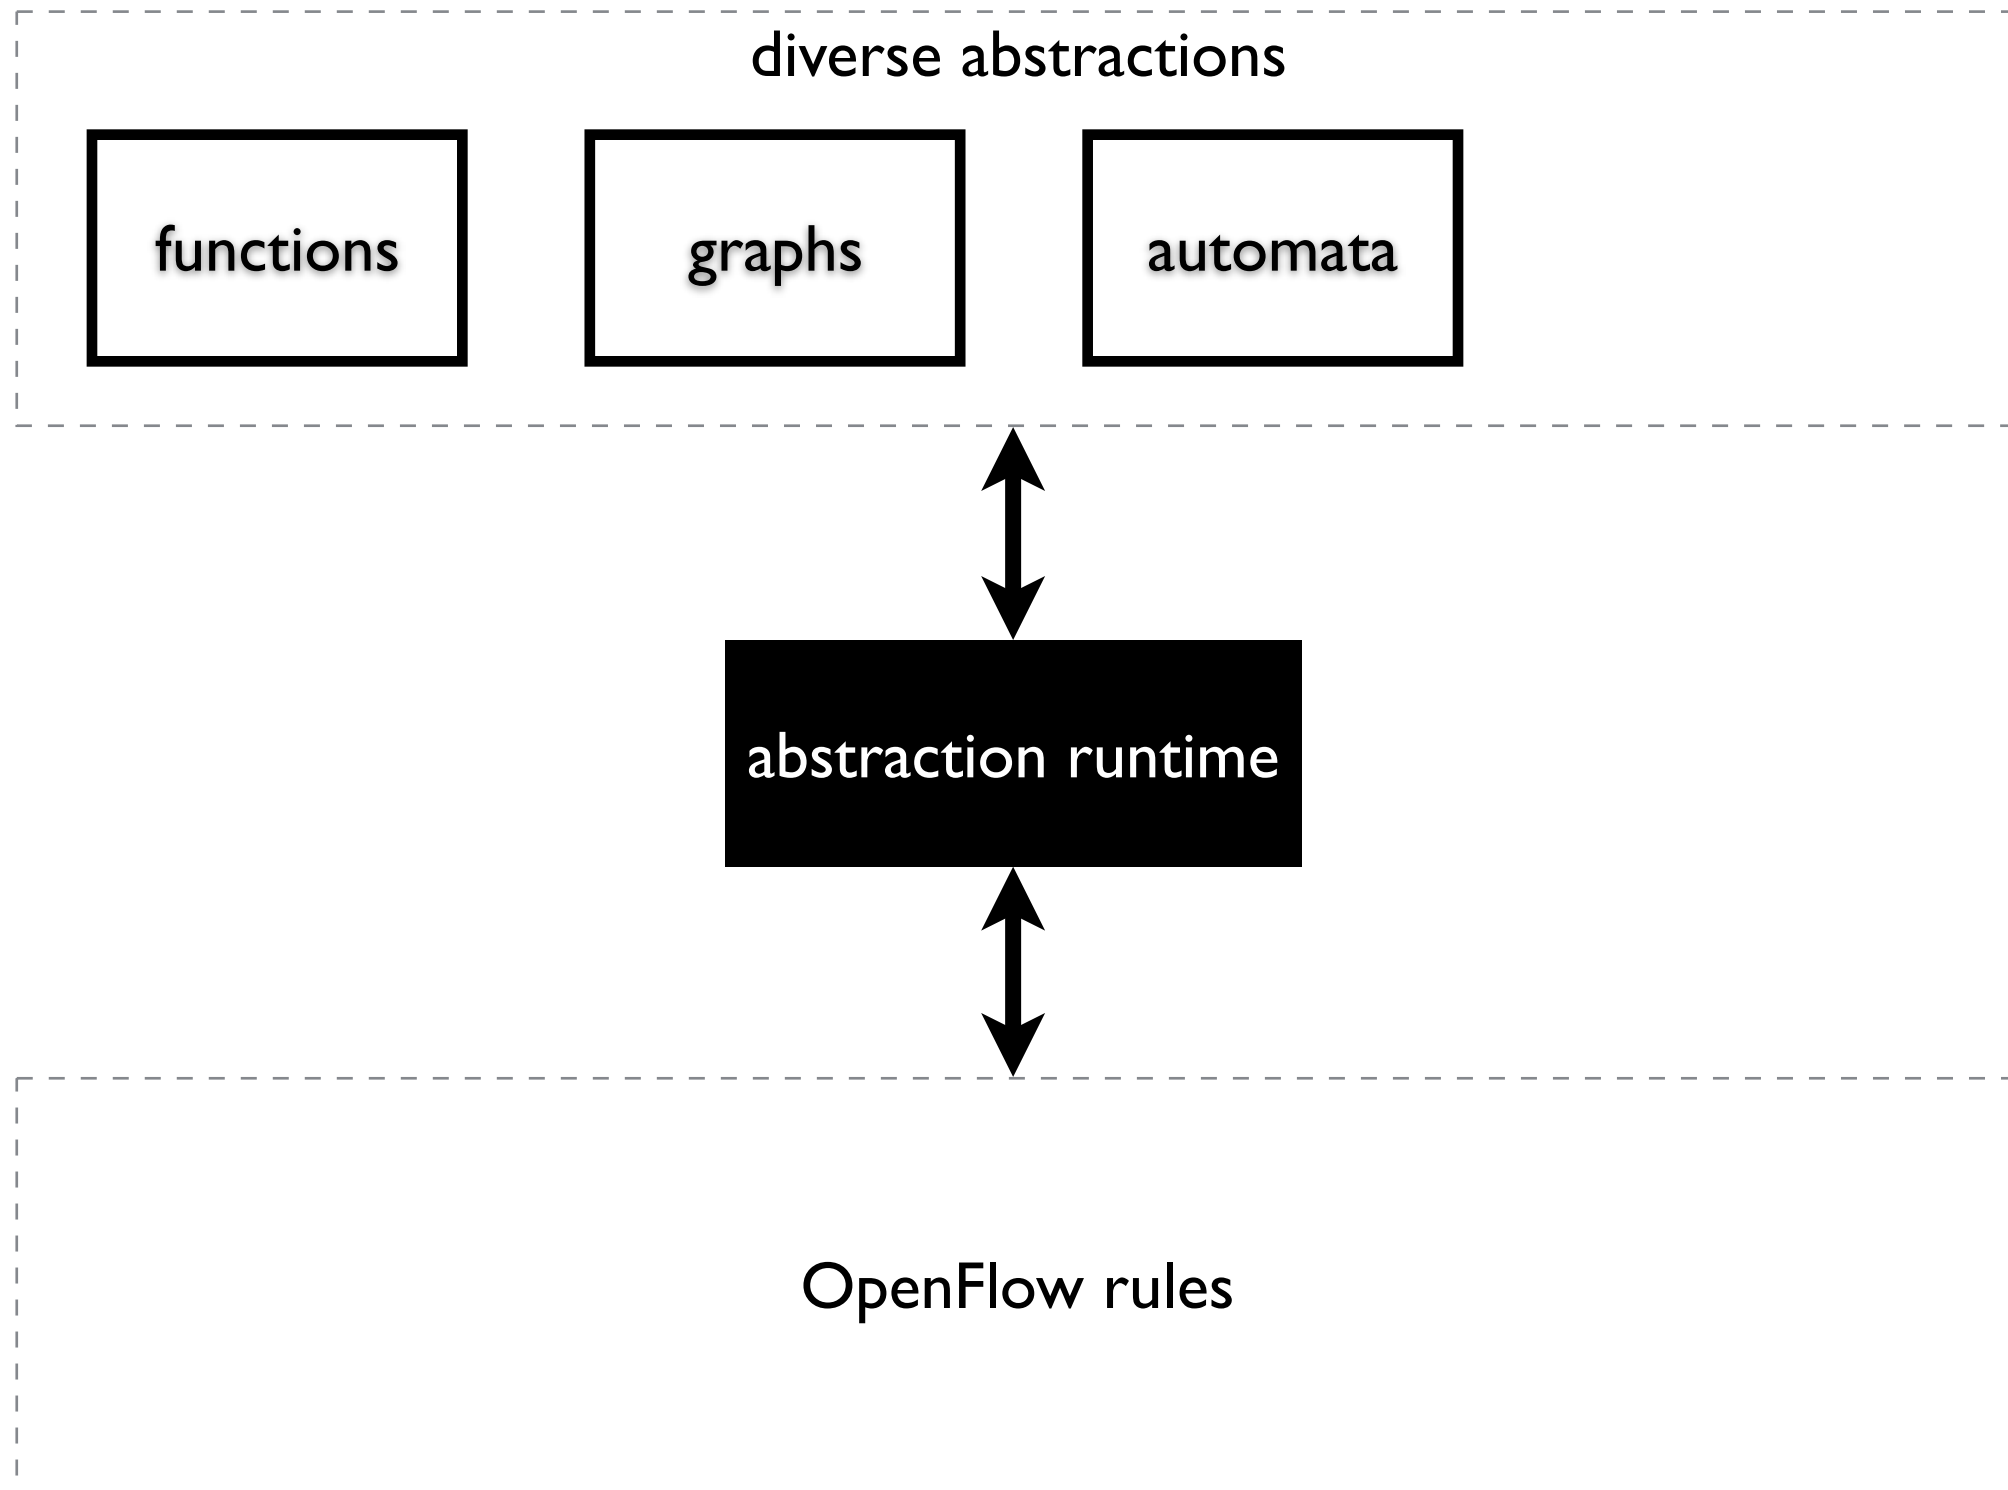
# abstractions



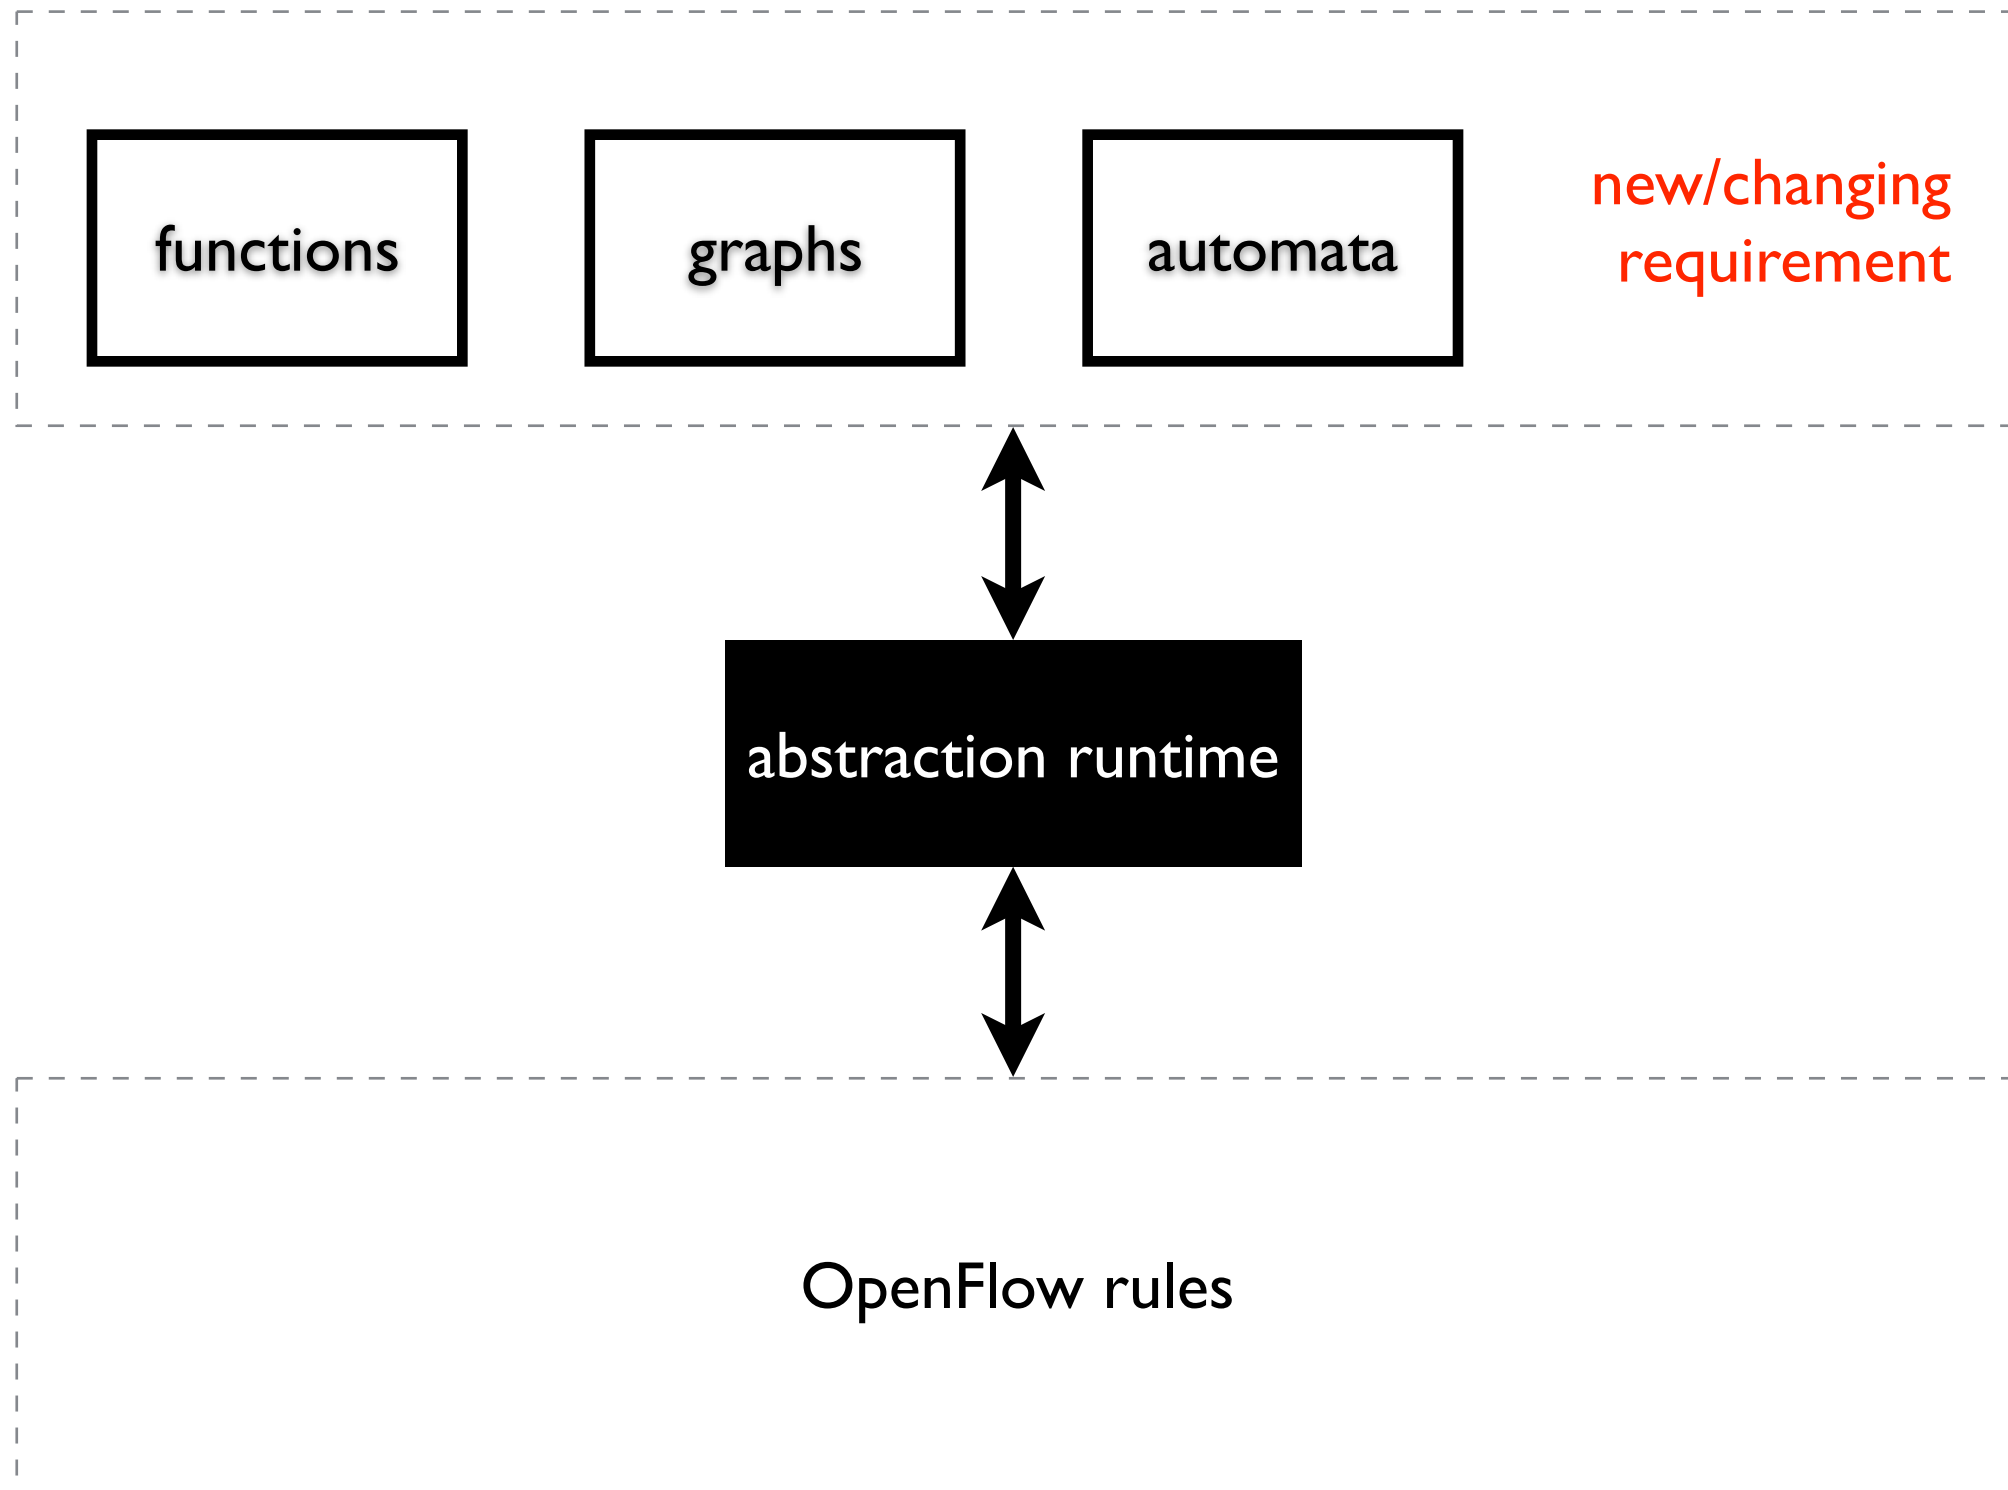
# abstractions



# abstractions

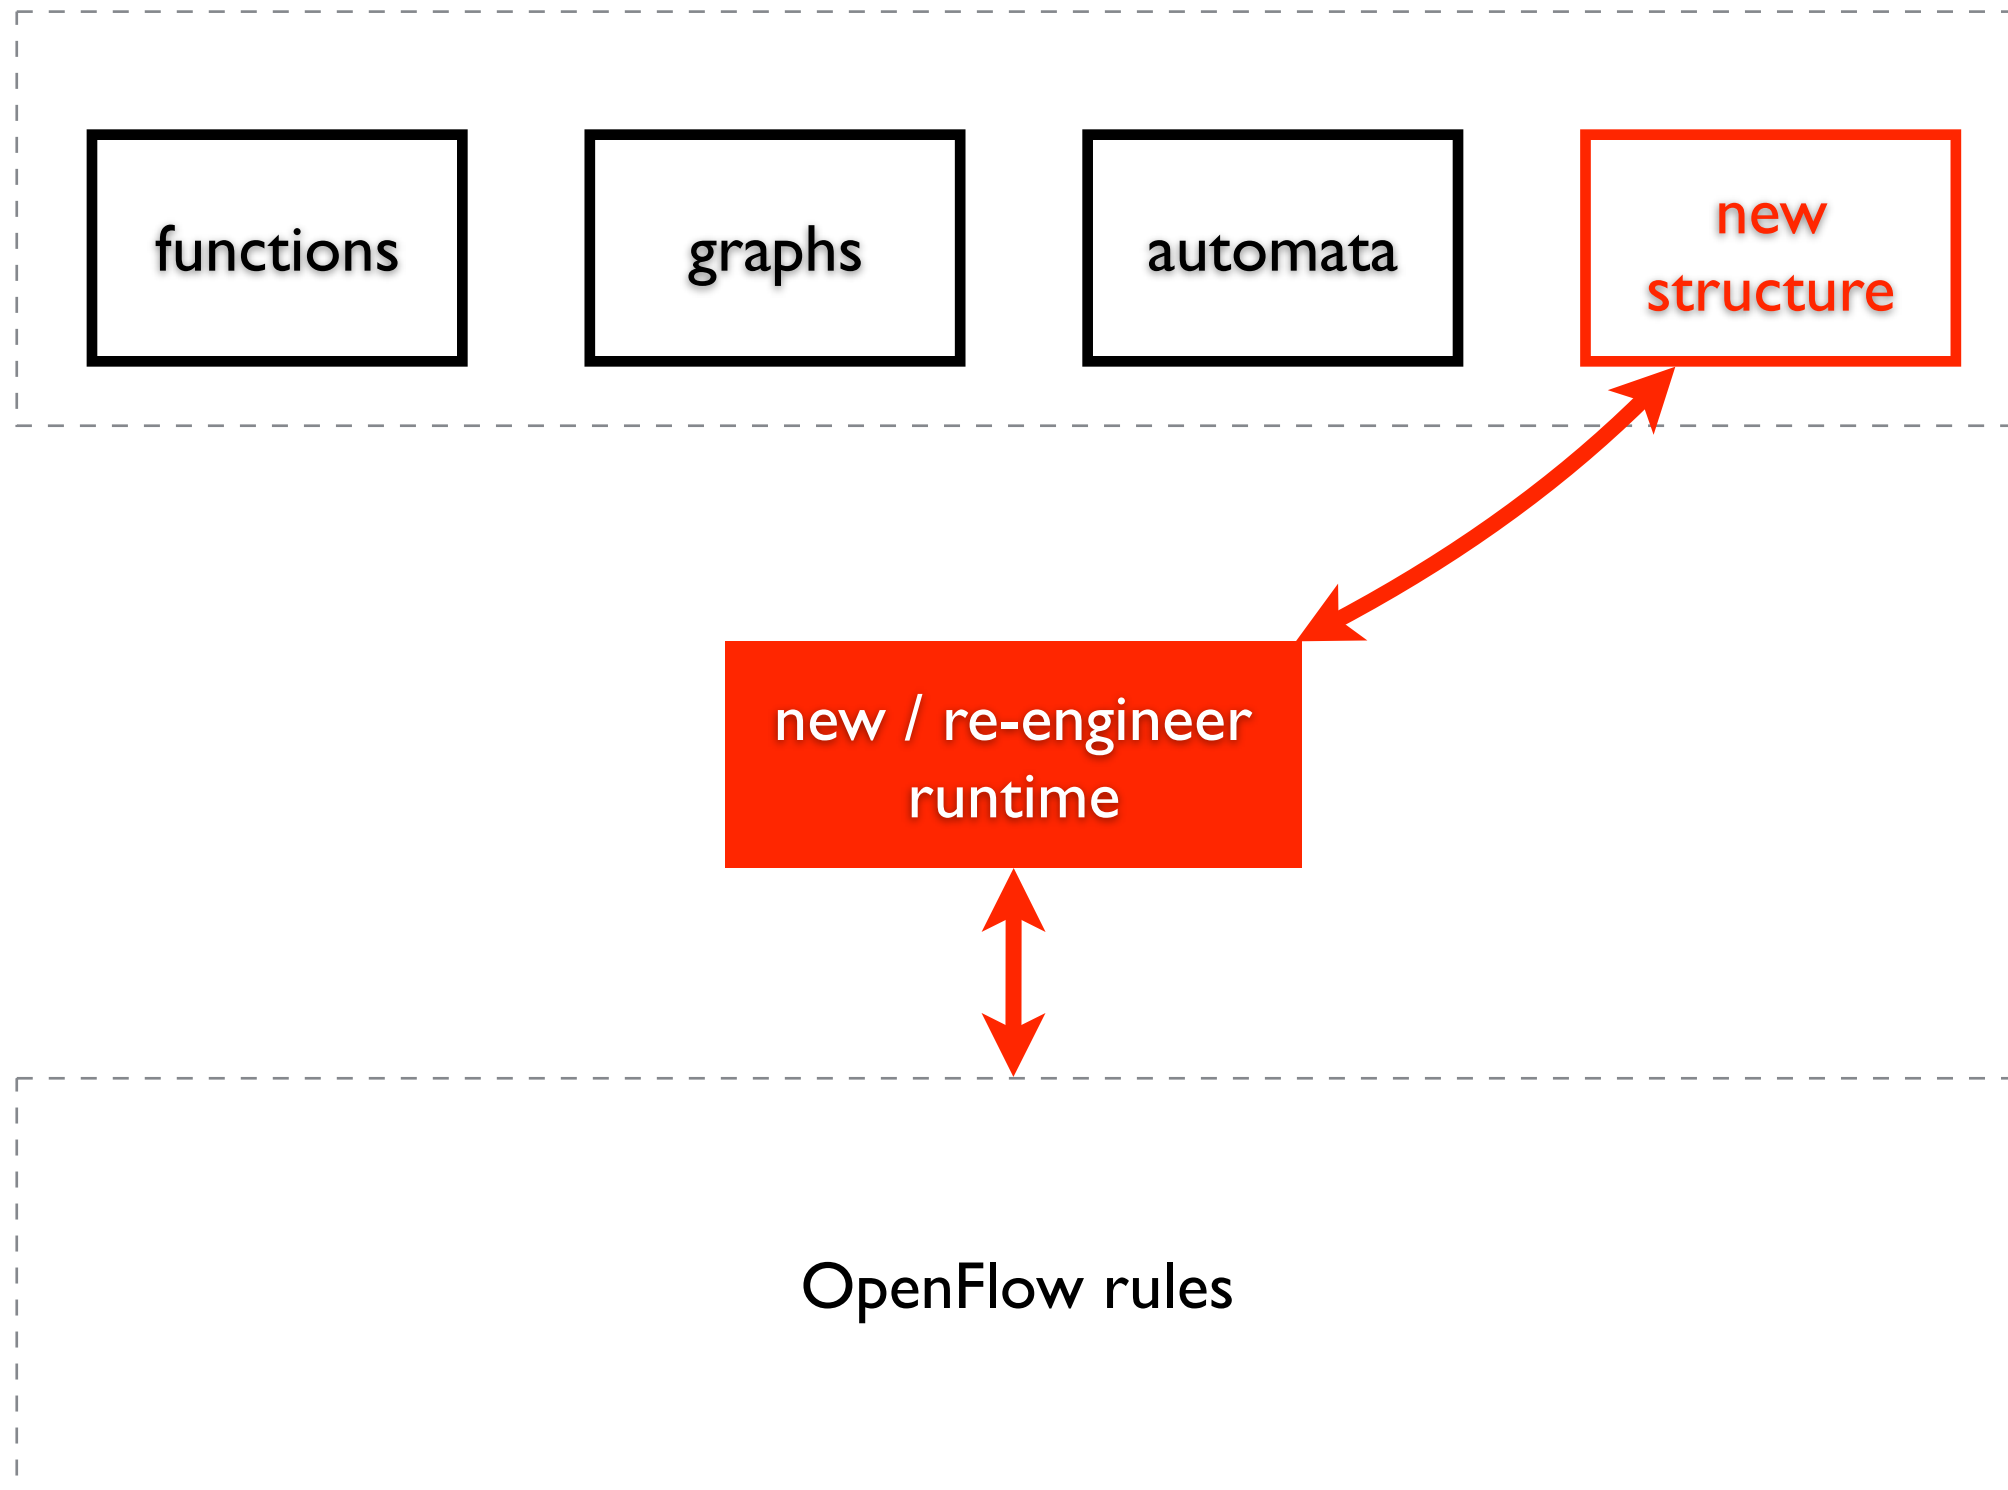


# but network keep evolving

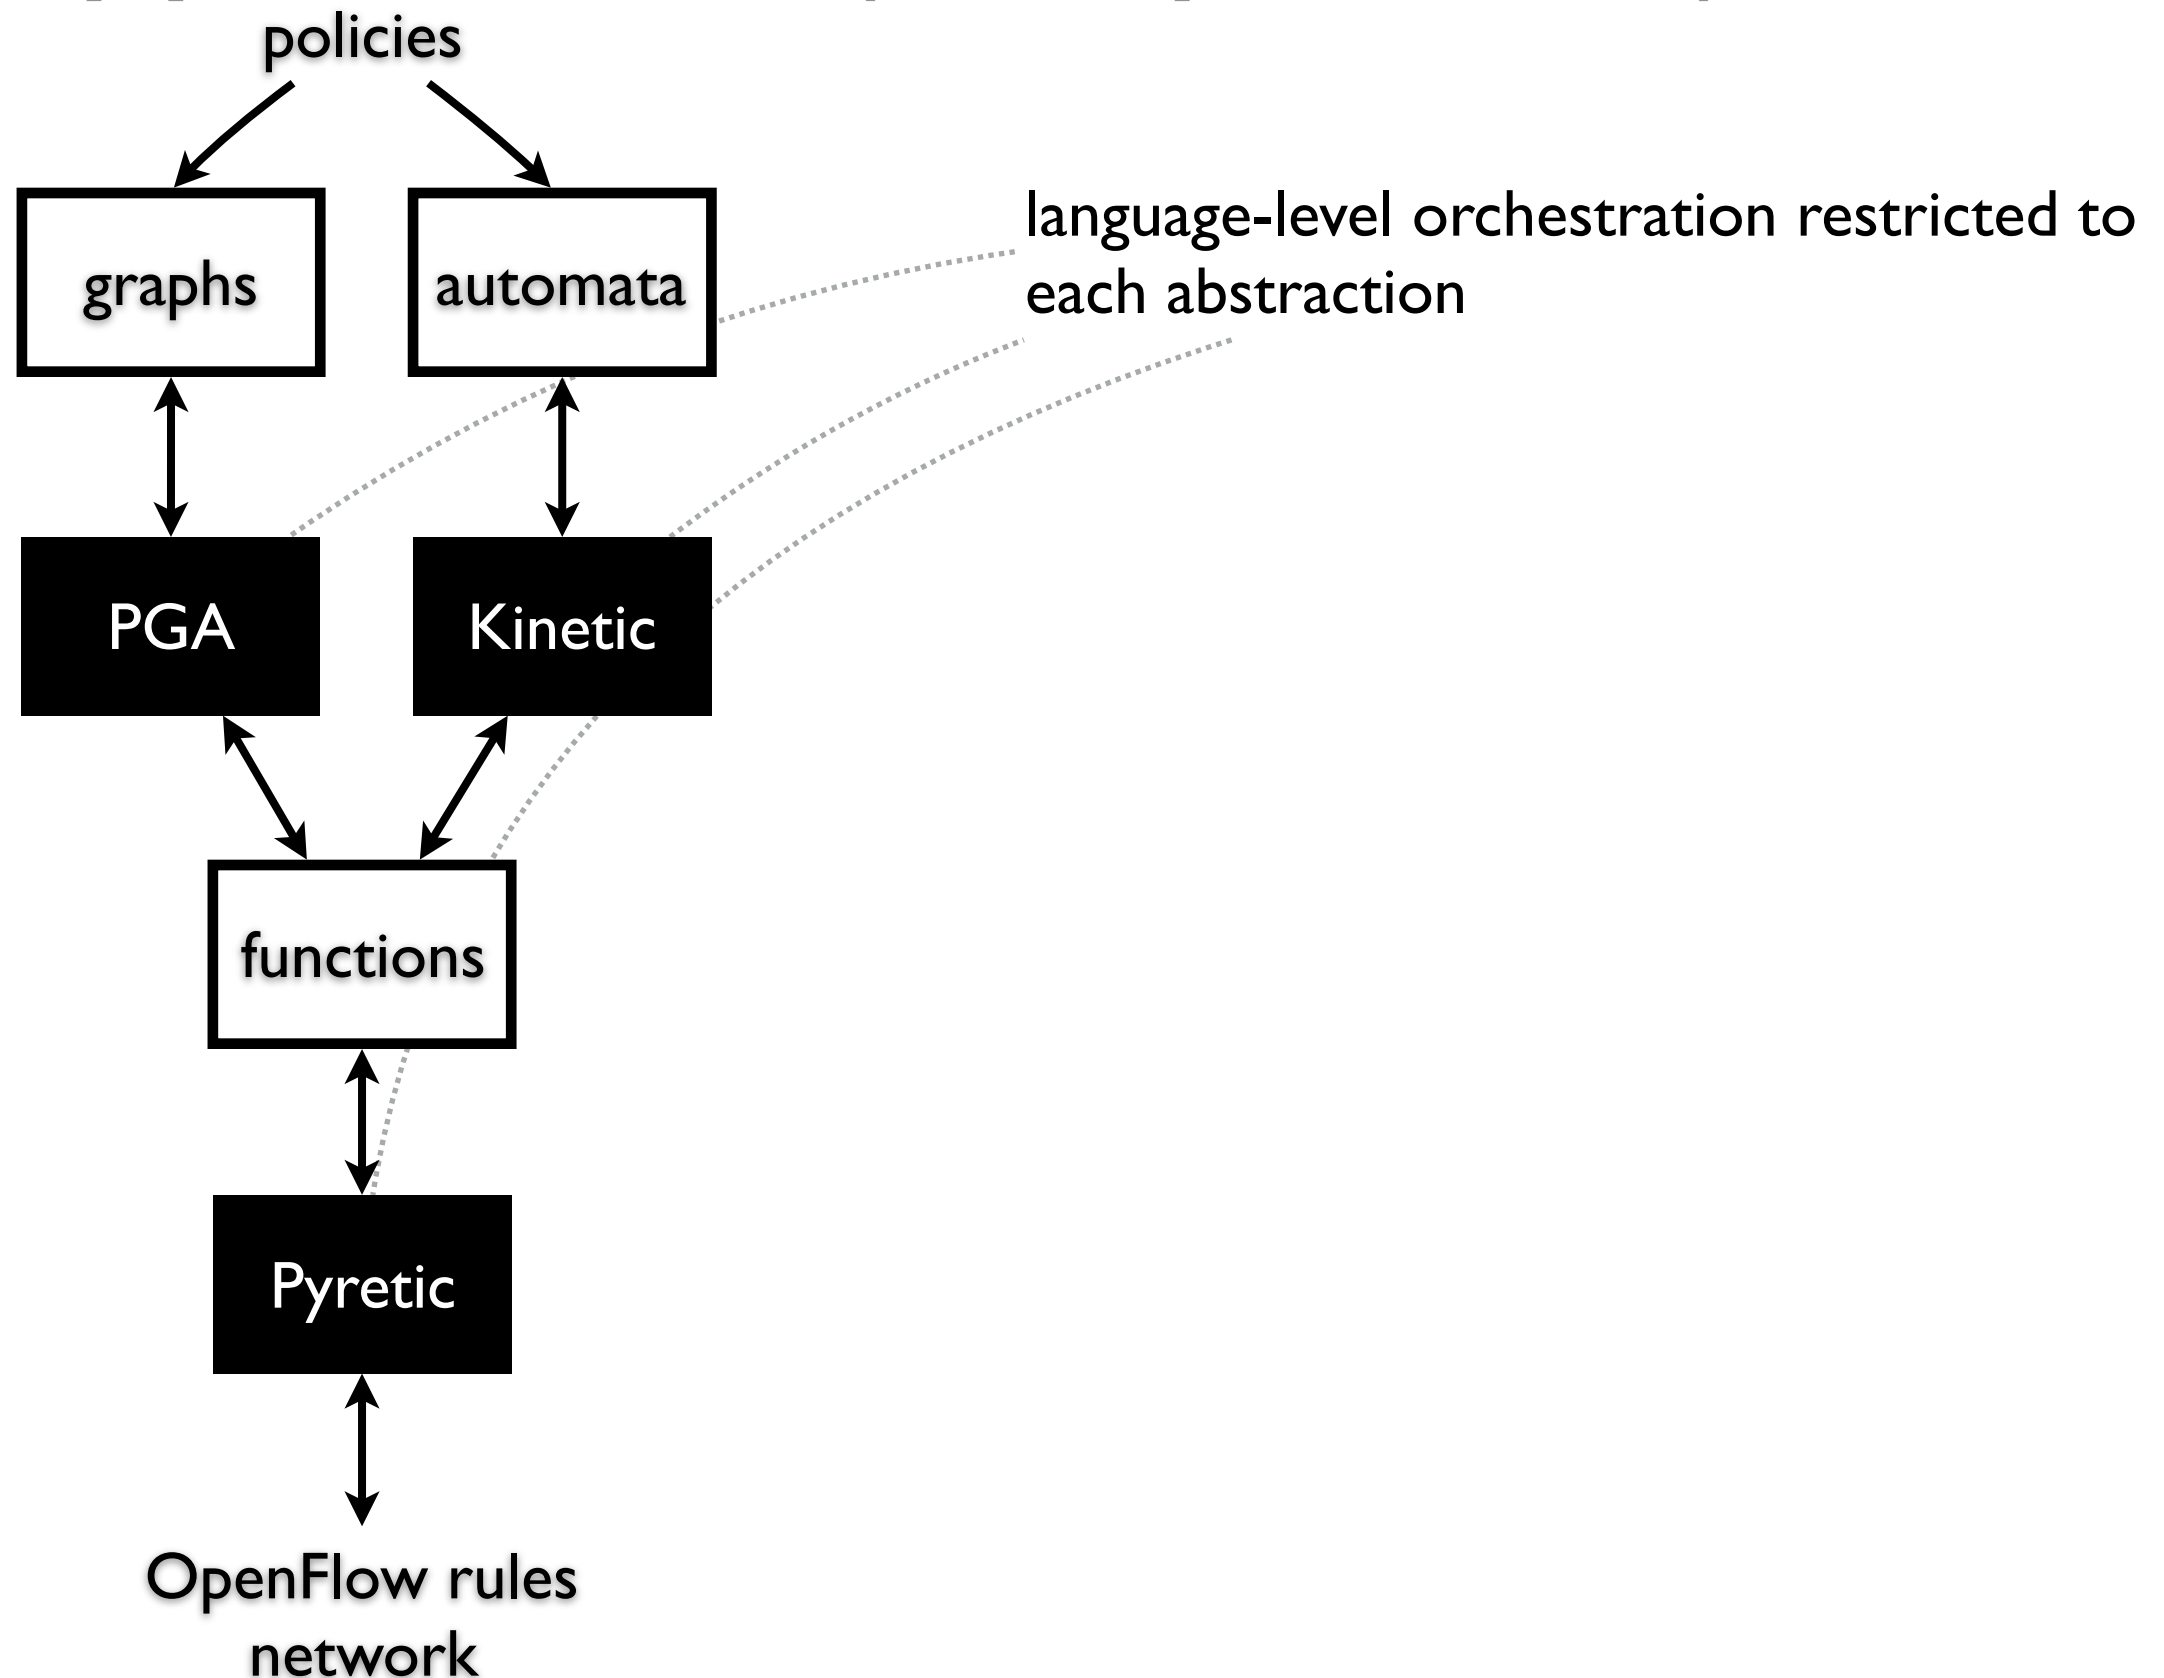




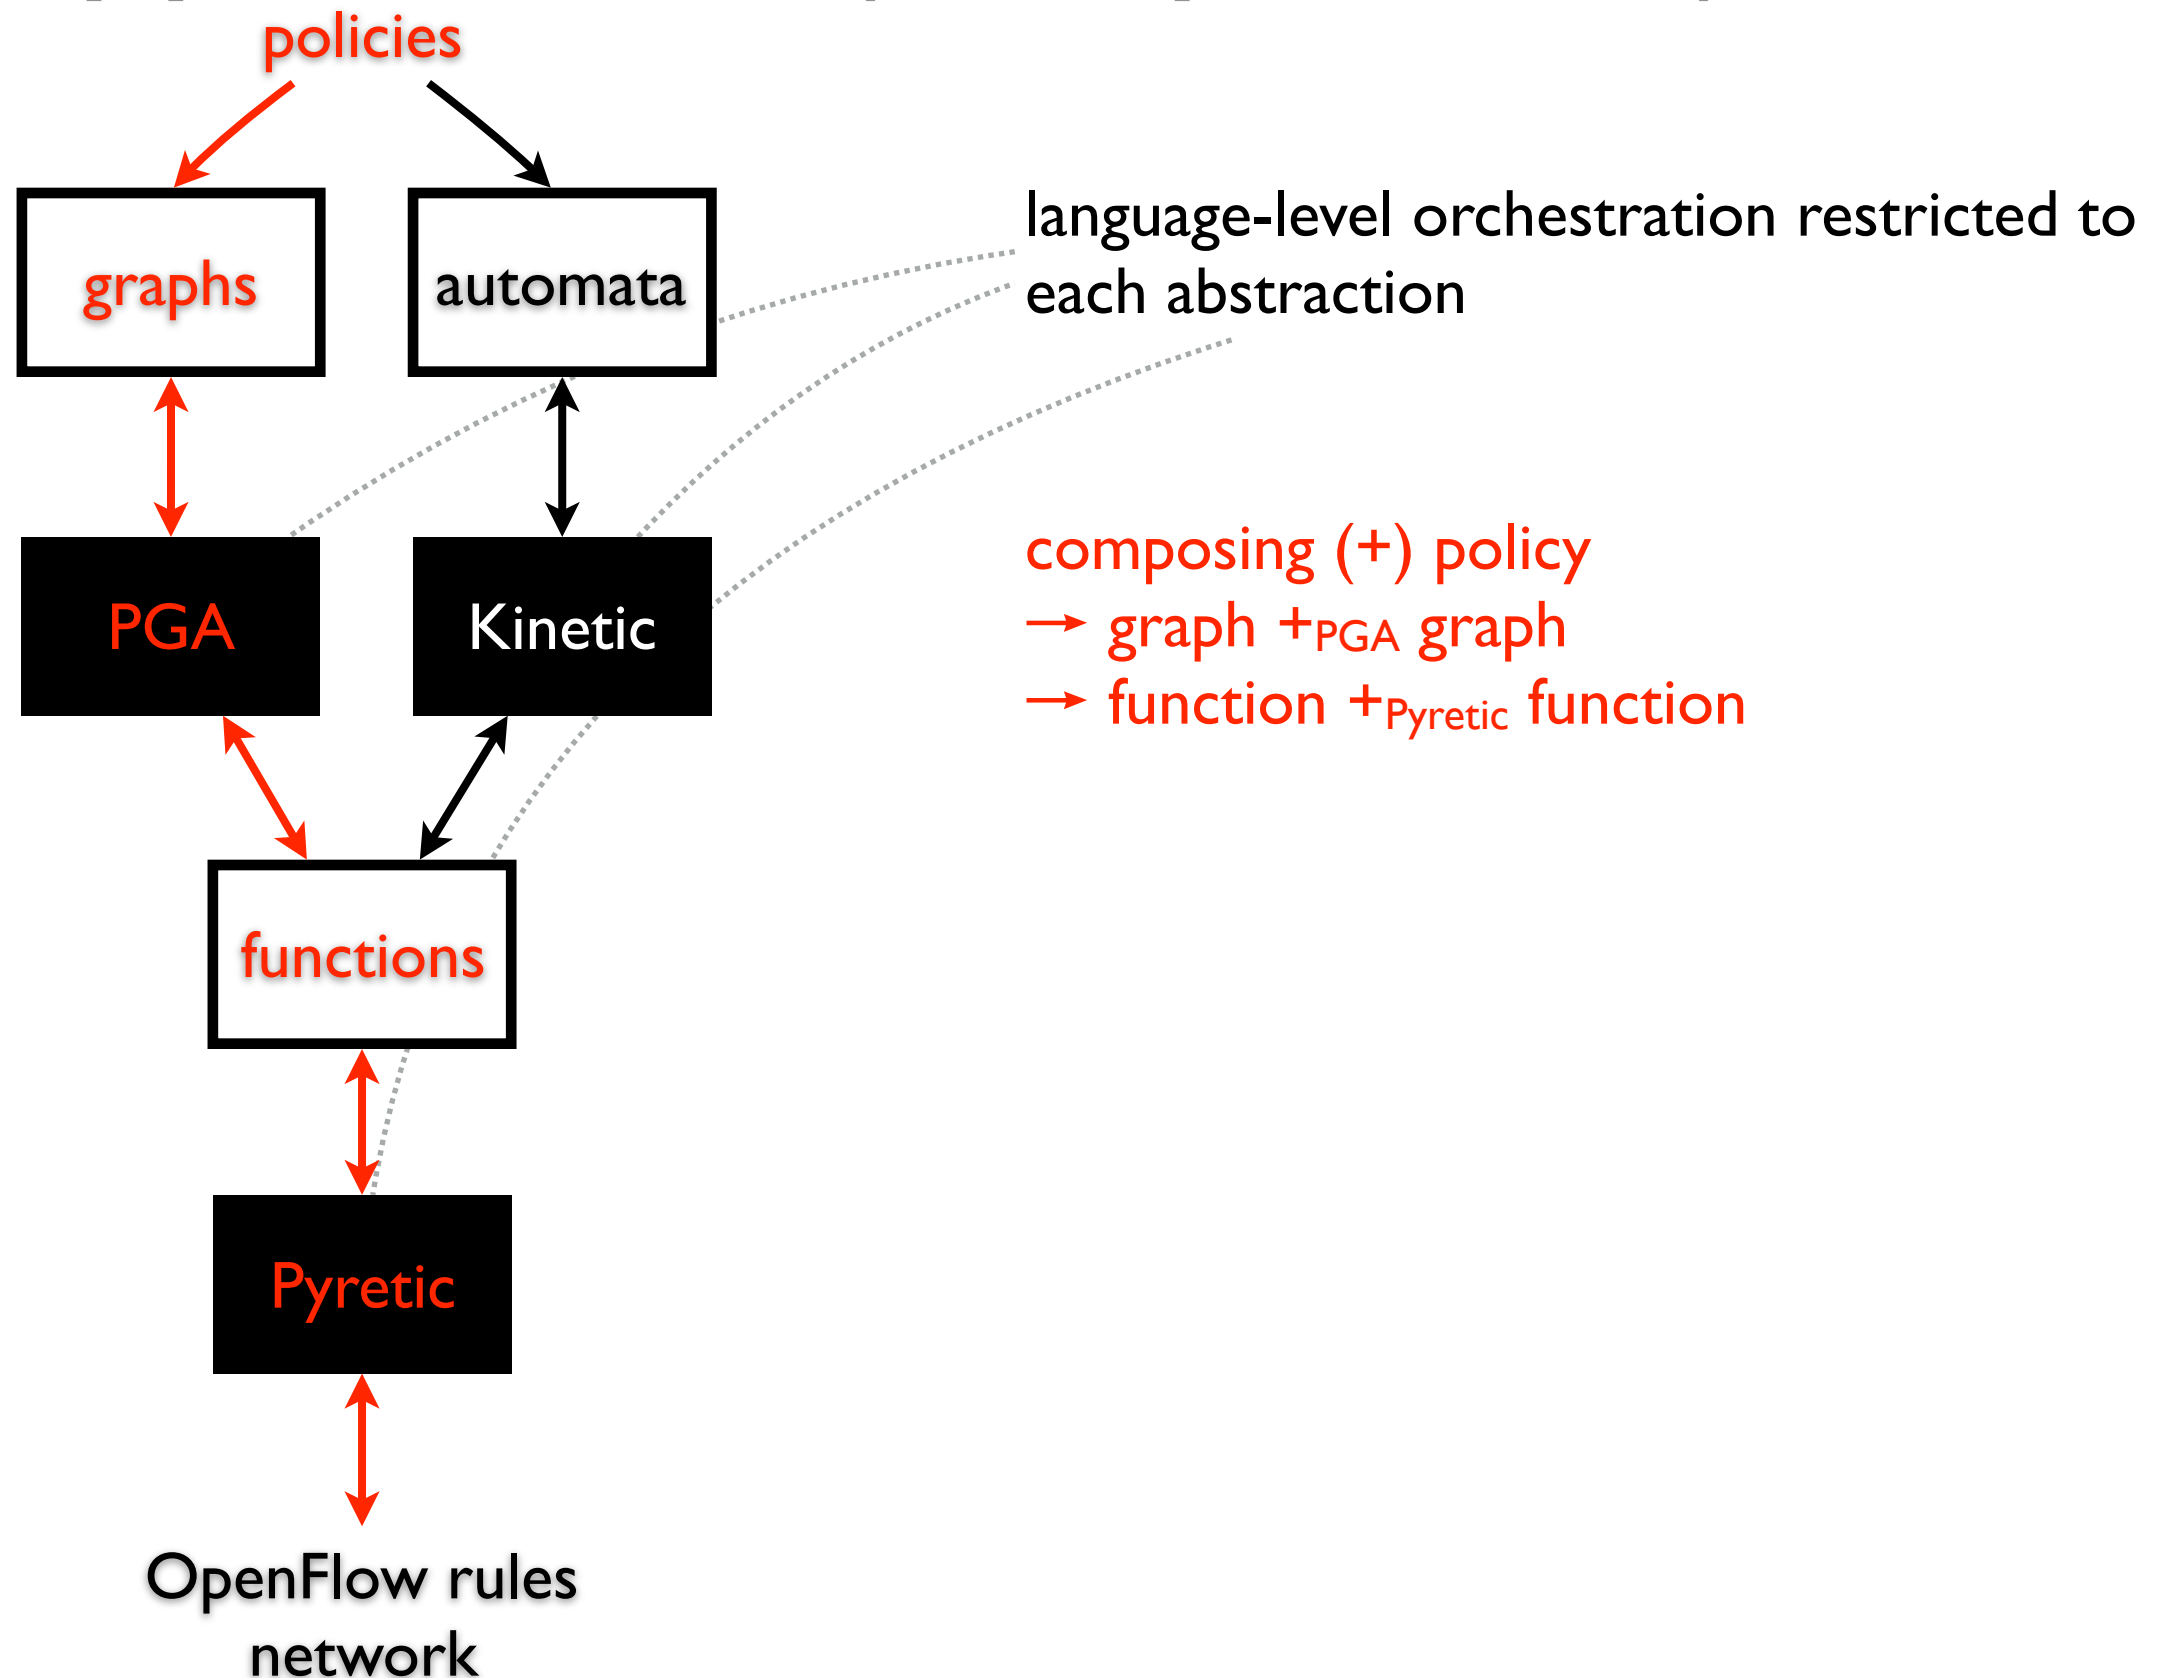
# but network keep evolving



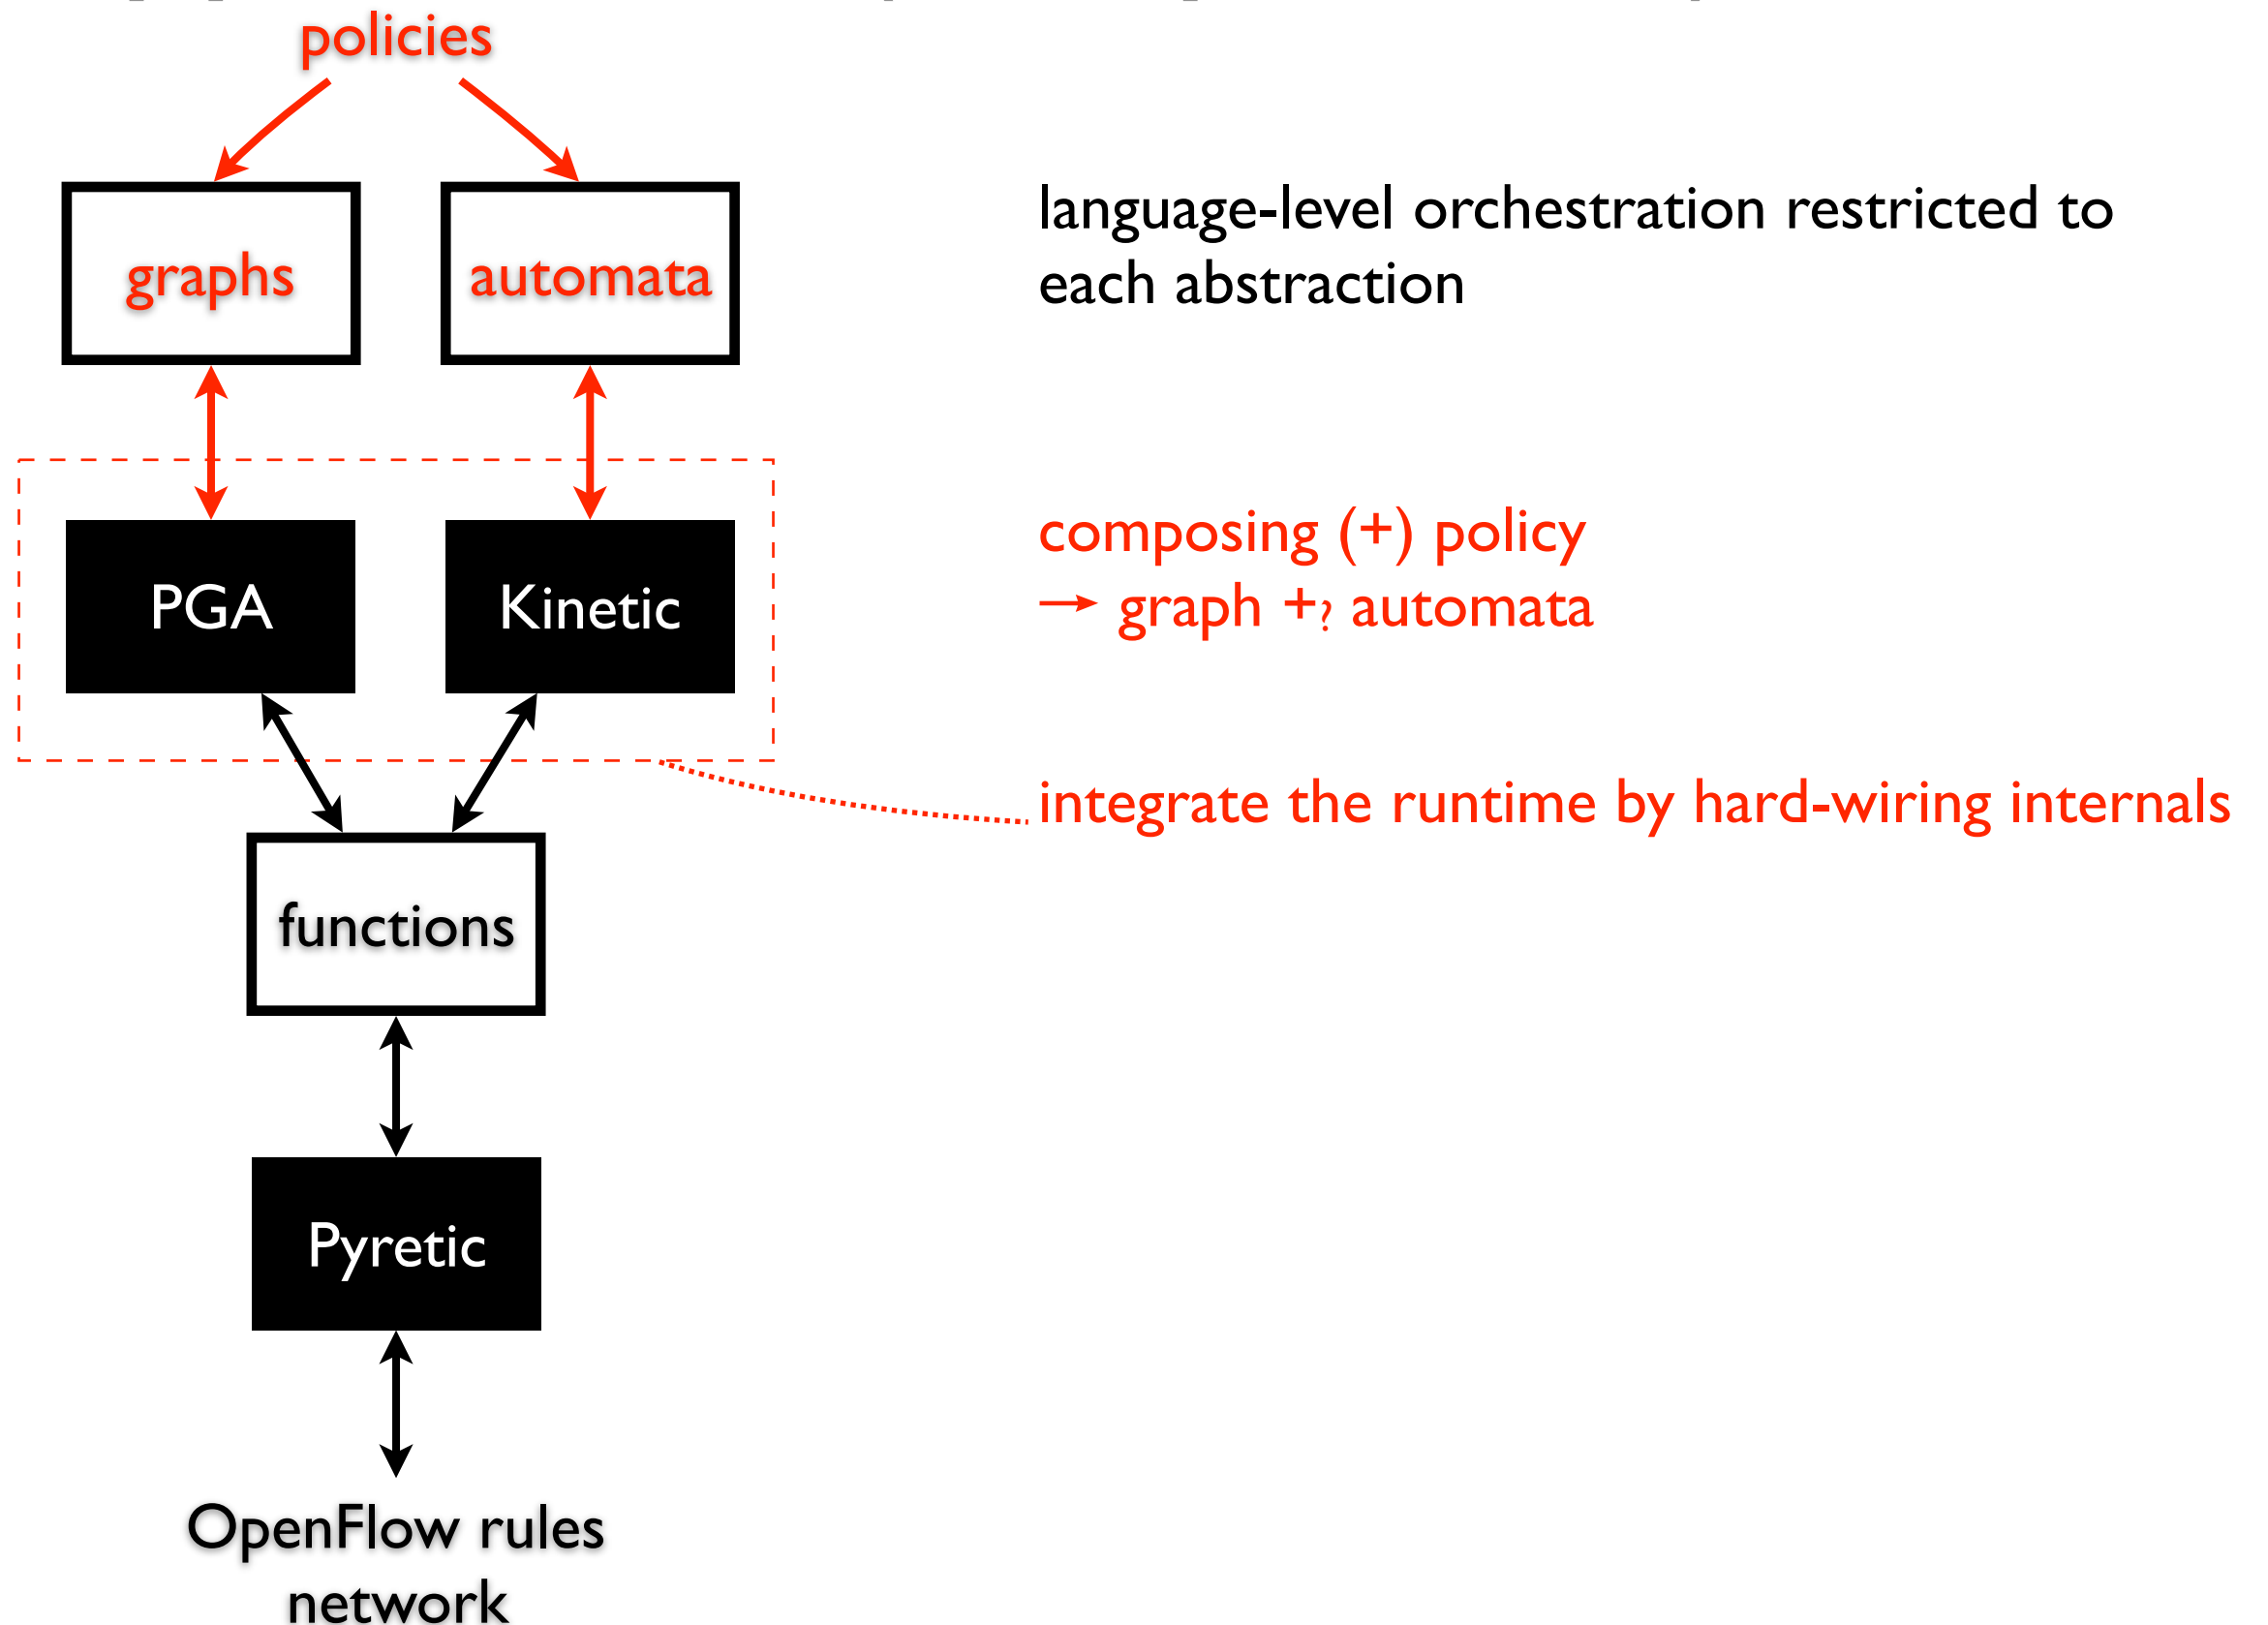
# and applications (components) interact



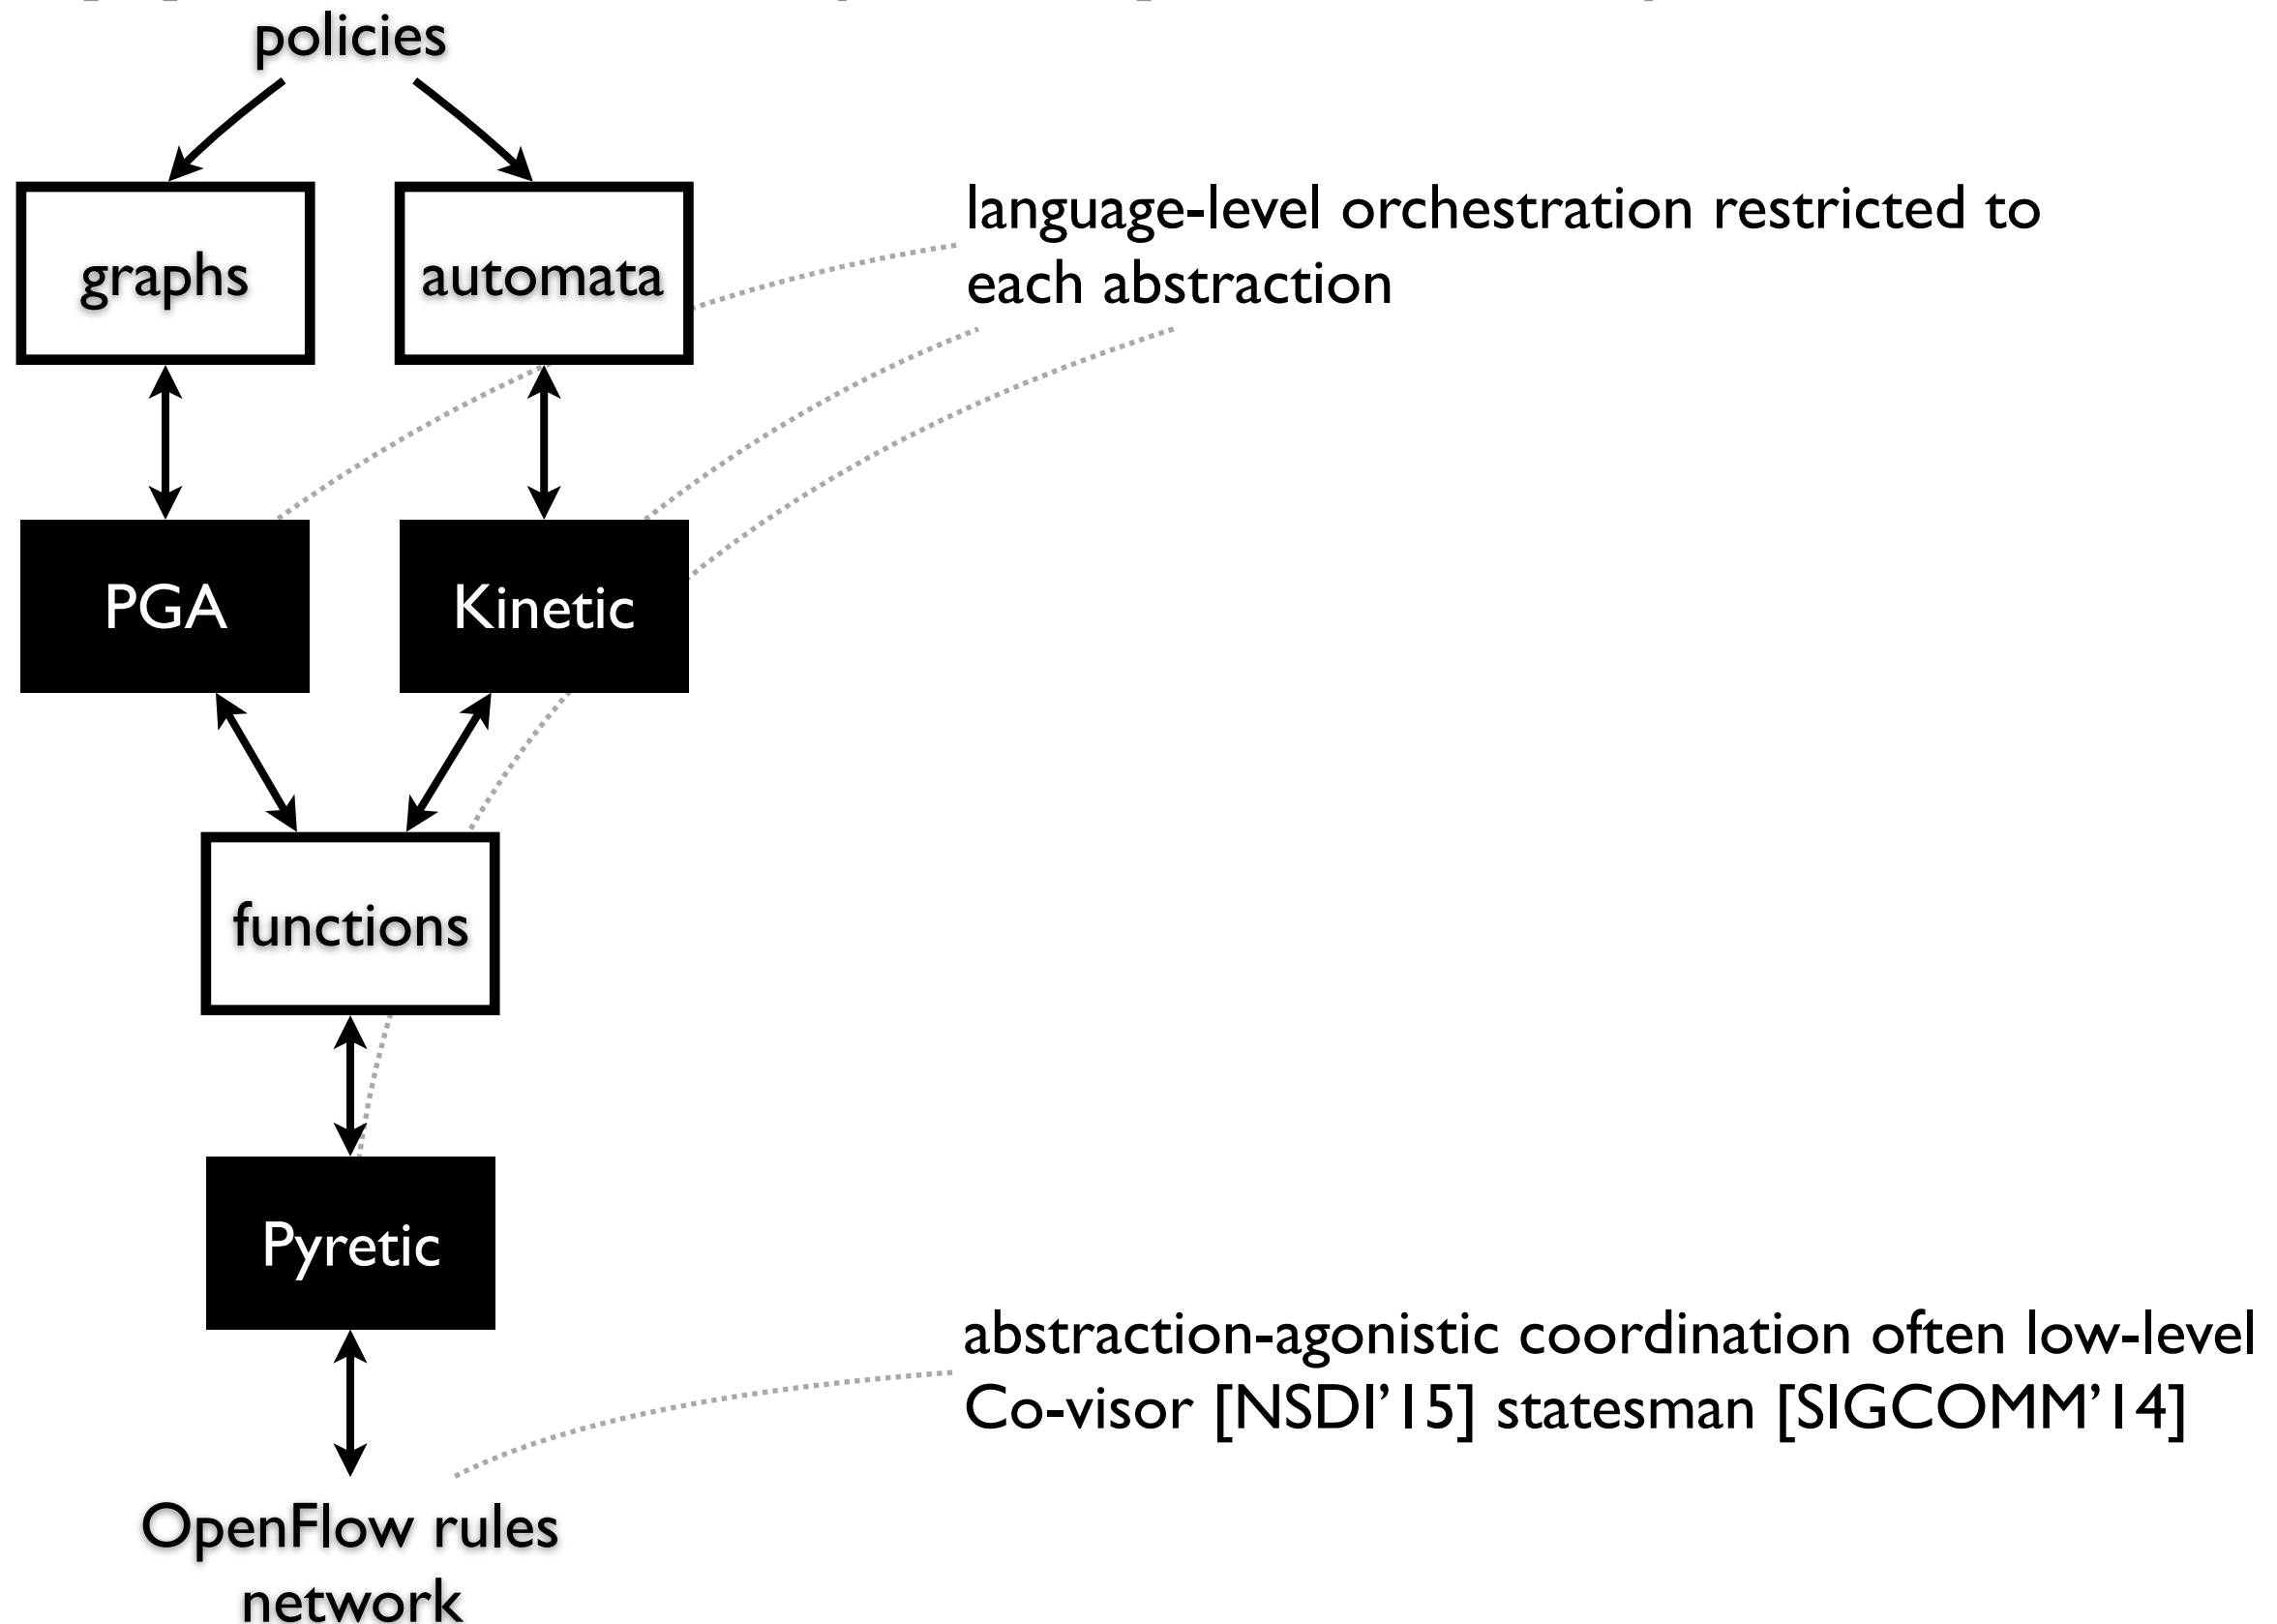
# and applications (components) interact



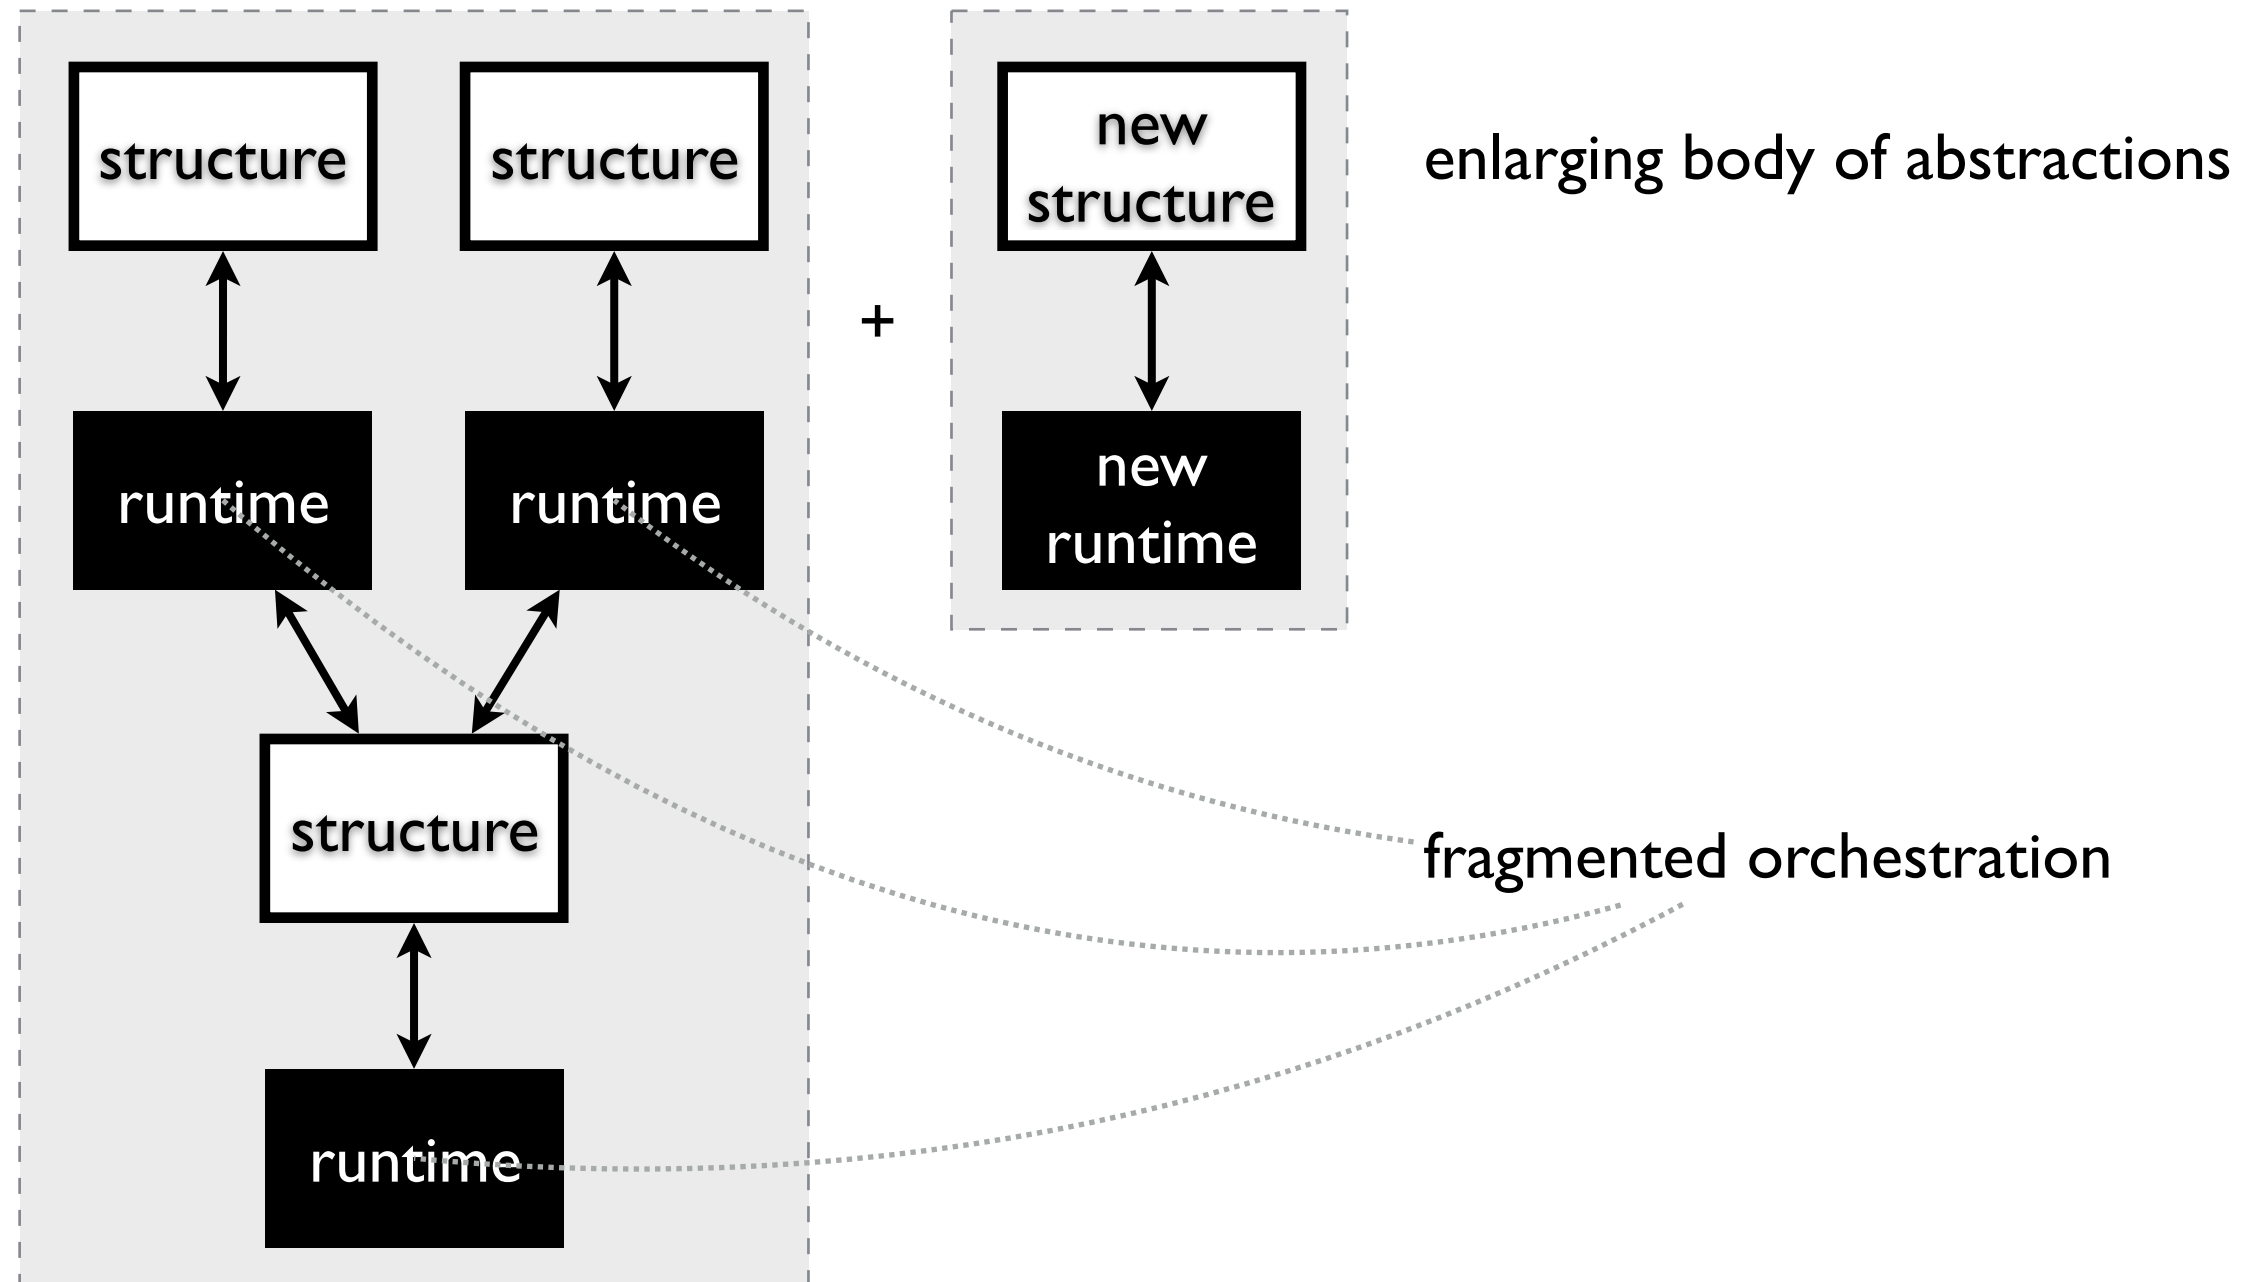
# and applications (components) interact



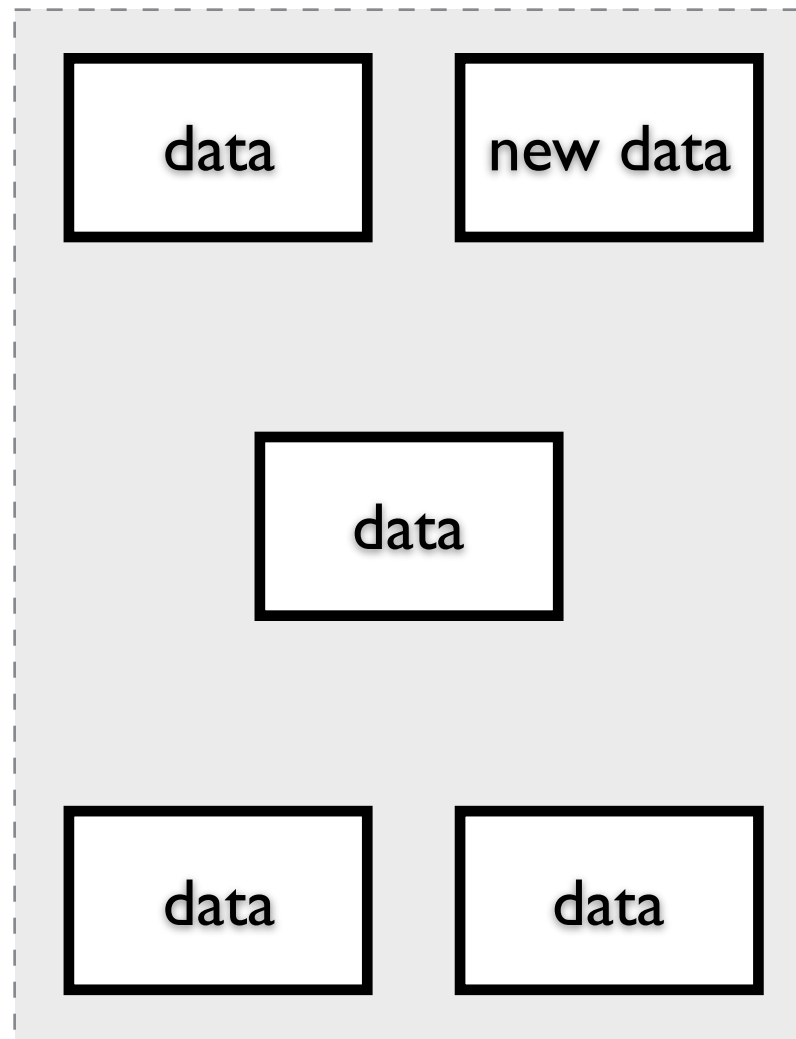
# and applications (components) interact



# current states of abstraction



# our perspective




SDN control revolves around data representation

- discard specialized, pre-compiled, fixed structure
- adopt a *plain data representation*

# our perspective

SDN control revolves around data representation

- discard specialized, pre-compiled, fixed structure
- adopt a *plain data representation*
- use a *universal data language*






# a database-defined network

operator and/or application

- **relation** — the plain data representation


- table — stored relation

- view — virtual relation




OpenFlow rules  
network

# a database-defined network




- **relation** — the plain data representation
  - table — stored relation
  - view — virtual relation
- **SQL** — the universal data language
  - SQL query ..... (dotted line)
  - SQL update — (solid line)
  - SQL trigger - - - - - (dashed line)

# a database-defined network



- relation — the plain data representation
  - table — stored relation
  - view — virtual relation
- SQL — the universal data language
  - SQL query .....
    - SQL update ———
    - SQL trigger ~~~~~
- SQL database — the high-performance runtime
  - orchestration challenge: refine runtime behavior by data mediation

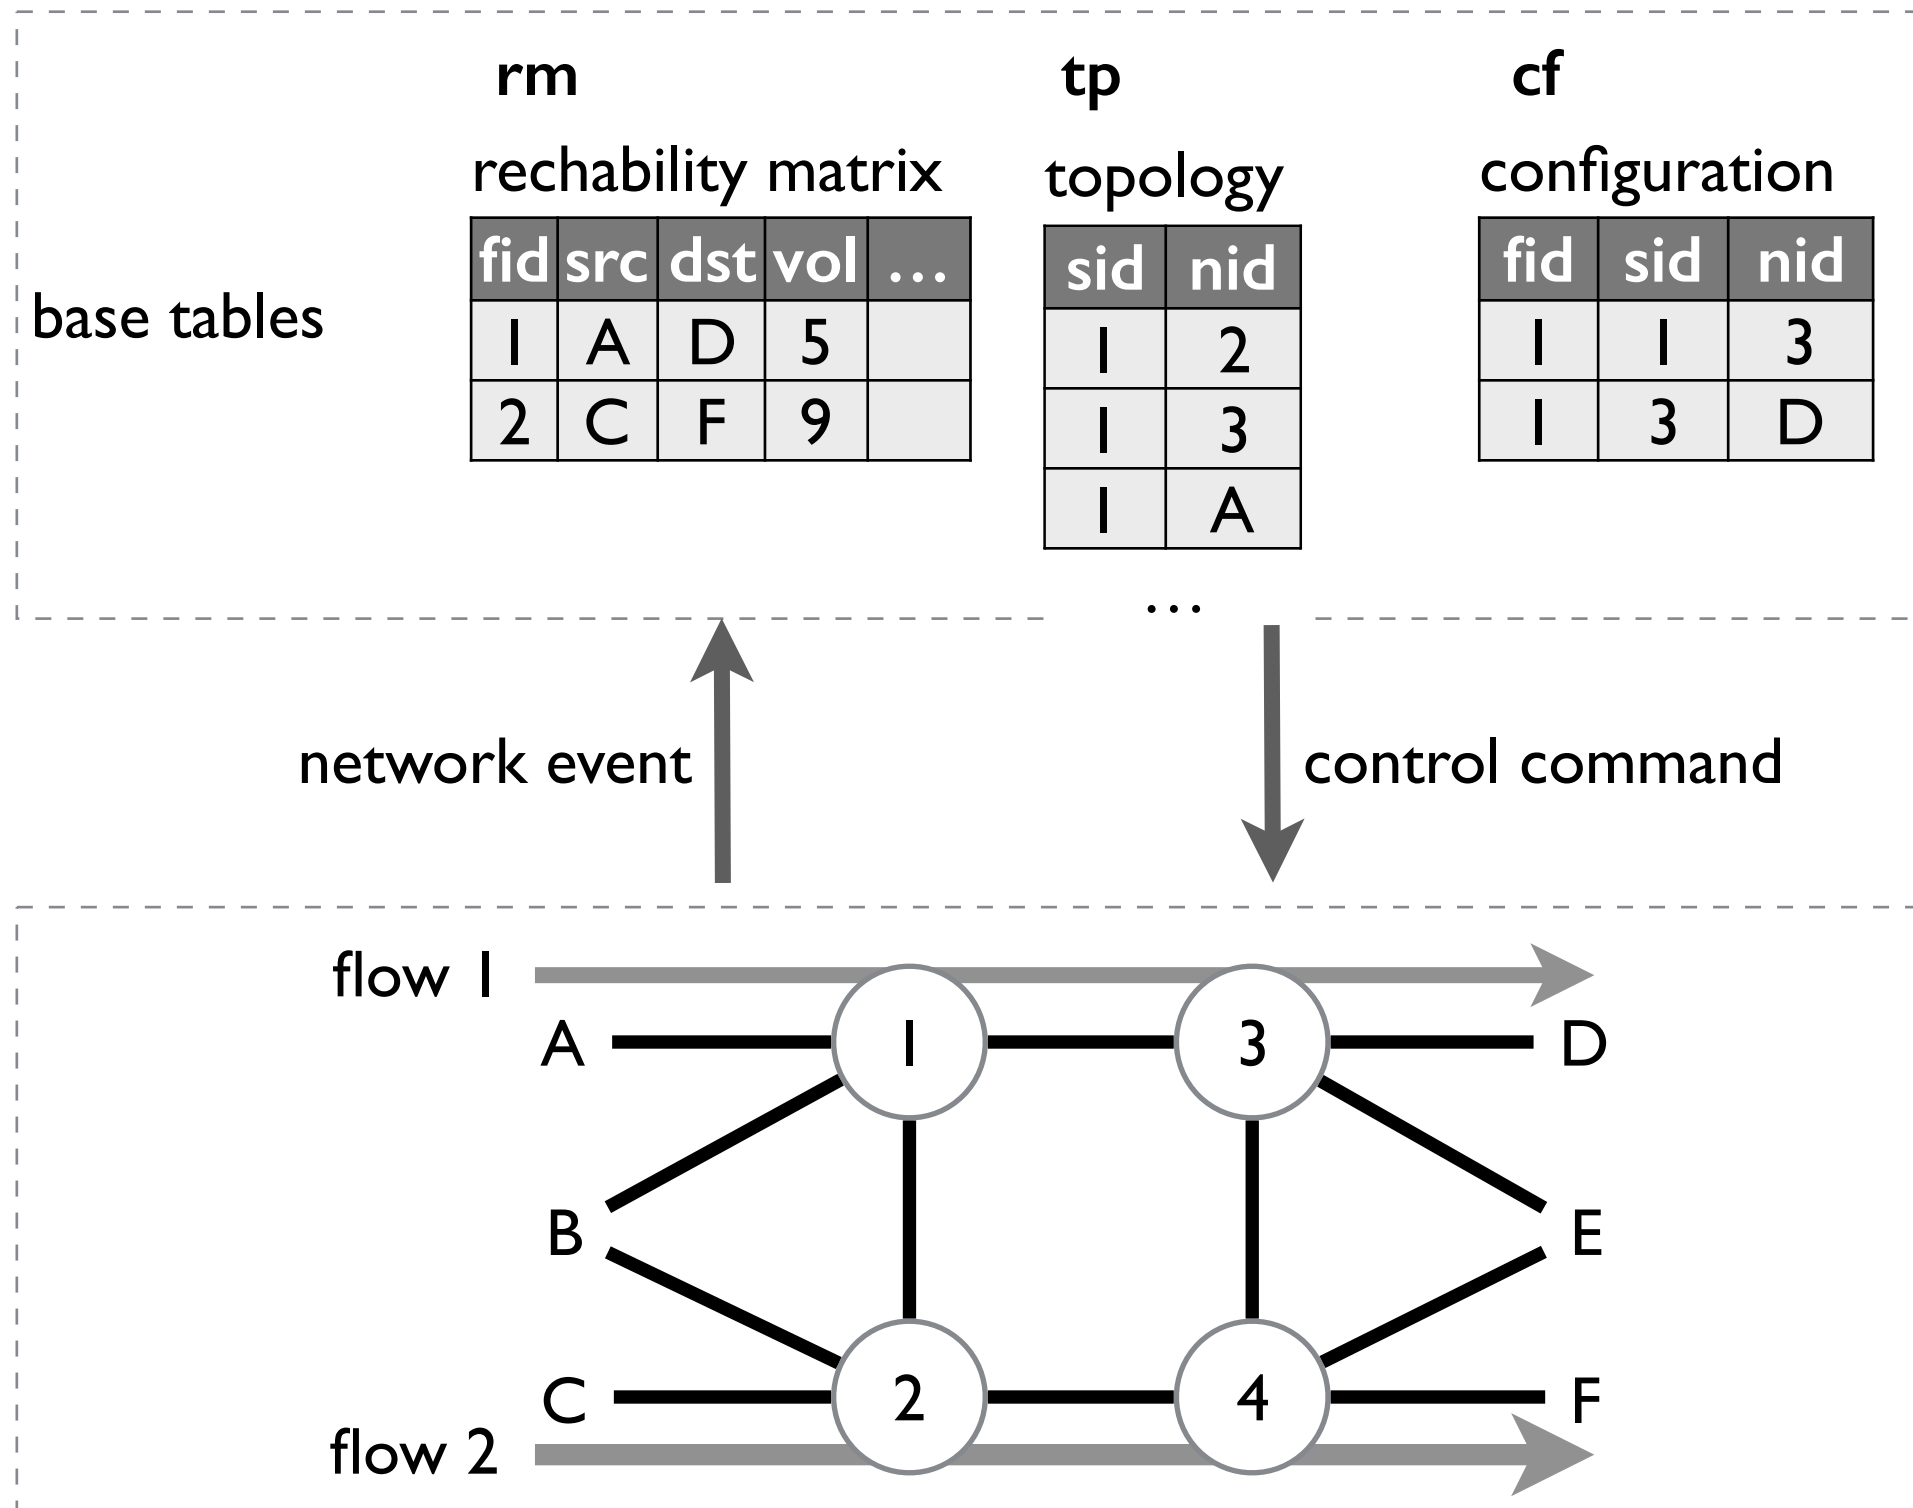
# Ravel: a realization with SQL database



## attractive features

- ad-hoc programmable abstraction via views
- orchestration across abstractions via view mechanism
- orchestration across applications via data mediation
- network control via SQL

# abstraction: network tables



# abstraction: application view

stateful firewall: allow inbound traffic only when initiated

firewall **tables**

FW: access control list

FW\_user: subnet

```
-- firewall state
CREATE TABLE FW (
  end1 integer, end2 integer, allow integer
);

CREATE TABLE FW_user (uid integer);
```

policy violation **view**

```
-- firewall constraints
CREATE OR REPLACE VIEW FW_violation AS (
  SELECT fid
  FROM tm
  WHERE FW = 1 AND
    (src, dst) NOT IN
    (SELECT end1, end2 FROM FW_policy_acl)
);
```

**rule** for violation repair

```
CREATE OR REPLACE RULE FW_repair AS
  ON DELETE TO FW_violation
  DO INSTEAD
    DELETE FROM tm WHERE fid = OLD.fid;
```

two firewall **rules**

that encode the state transitions

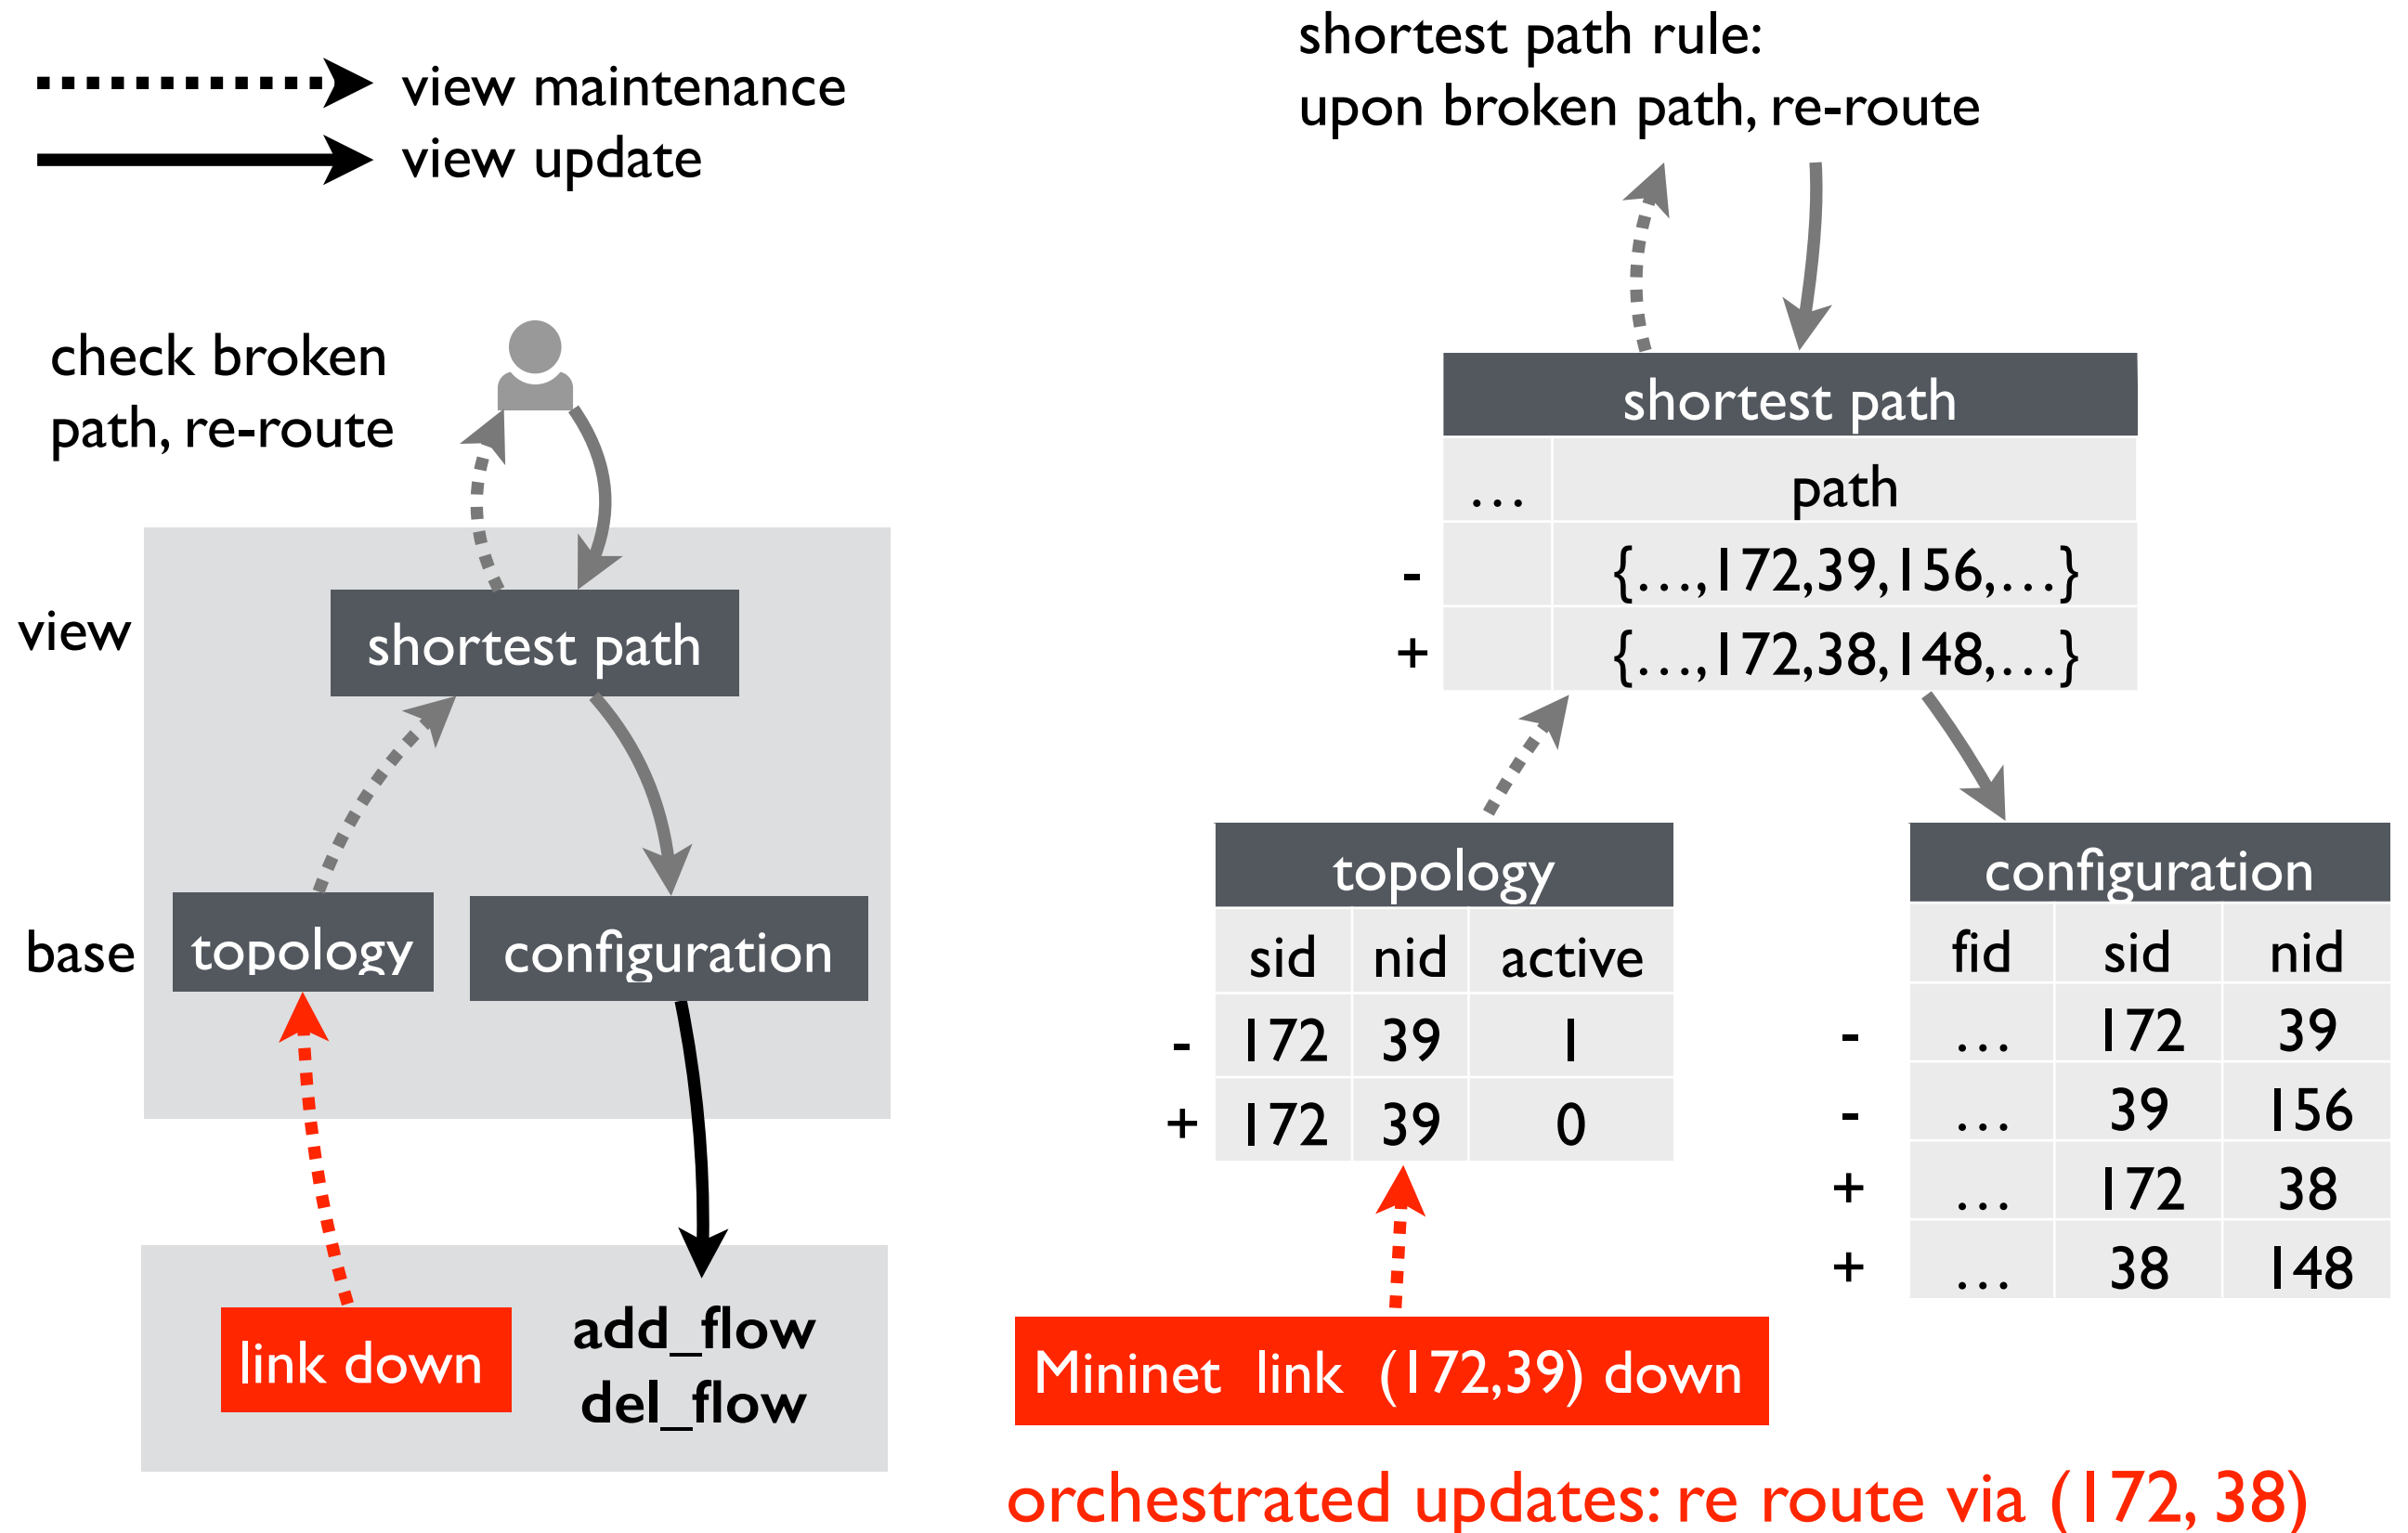
add allowed inbound routes

```
-- firewall state transitions
CREATE RULE FW1 AS
  ON INSERT TO tm
  WHERE ((NEW.src, NEW.dst) NOT IN
    (SELECT end2, end1 FROM FW)) AND
    (NEW.src IN (SELECT * FROM FW_user))
  DO ALSO (
    INSERT INTO FW
    VALUES (NEW.dst, NEW.src, 1));
```

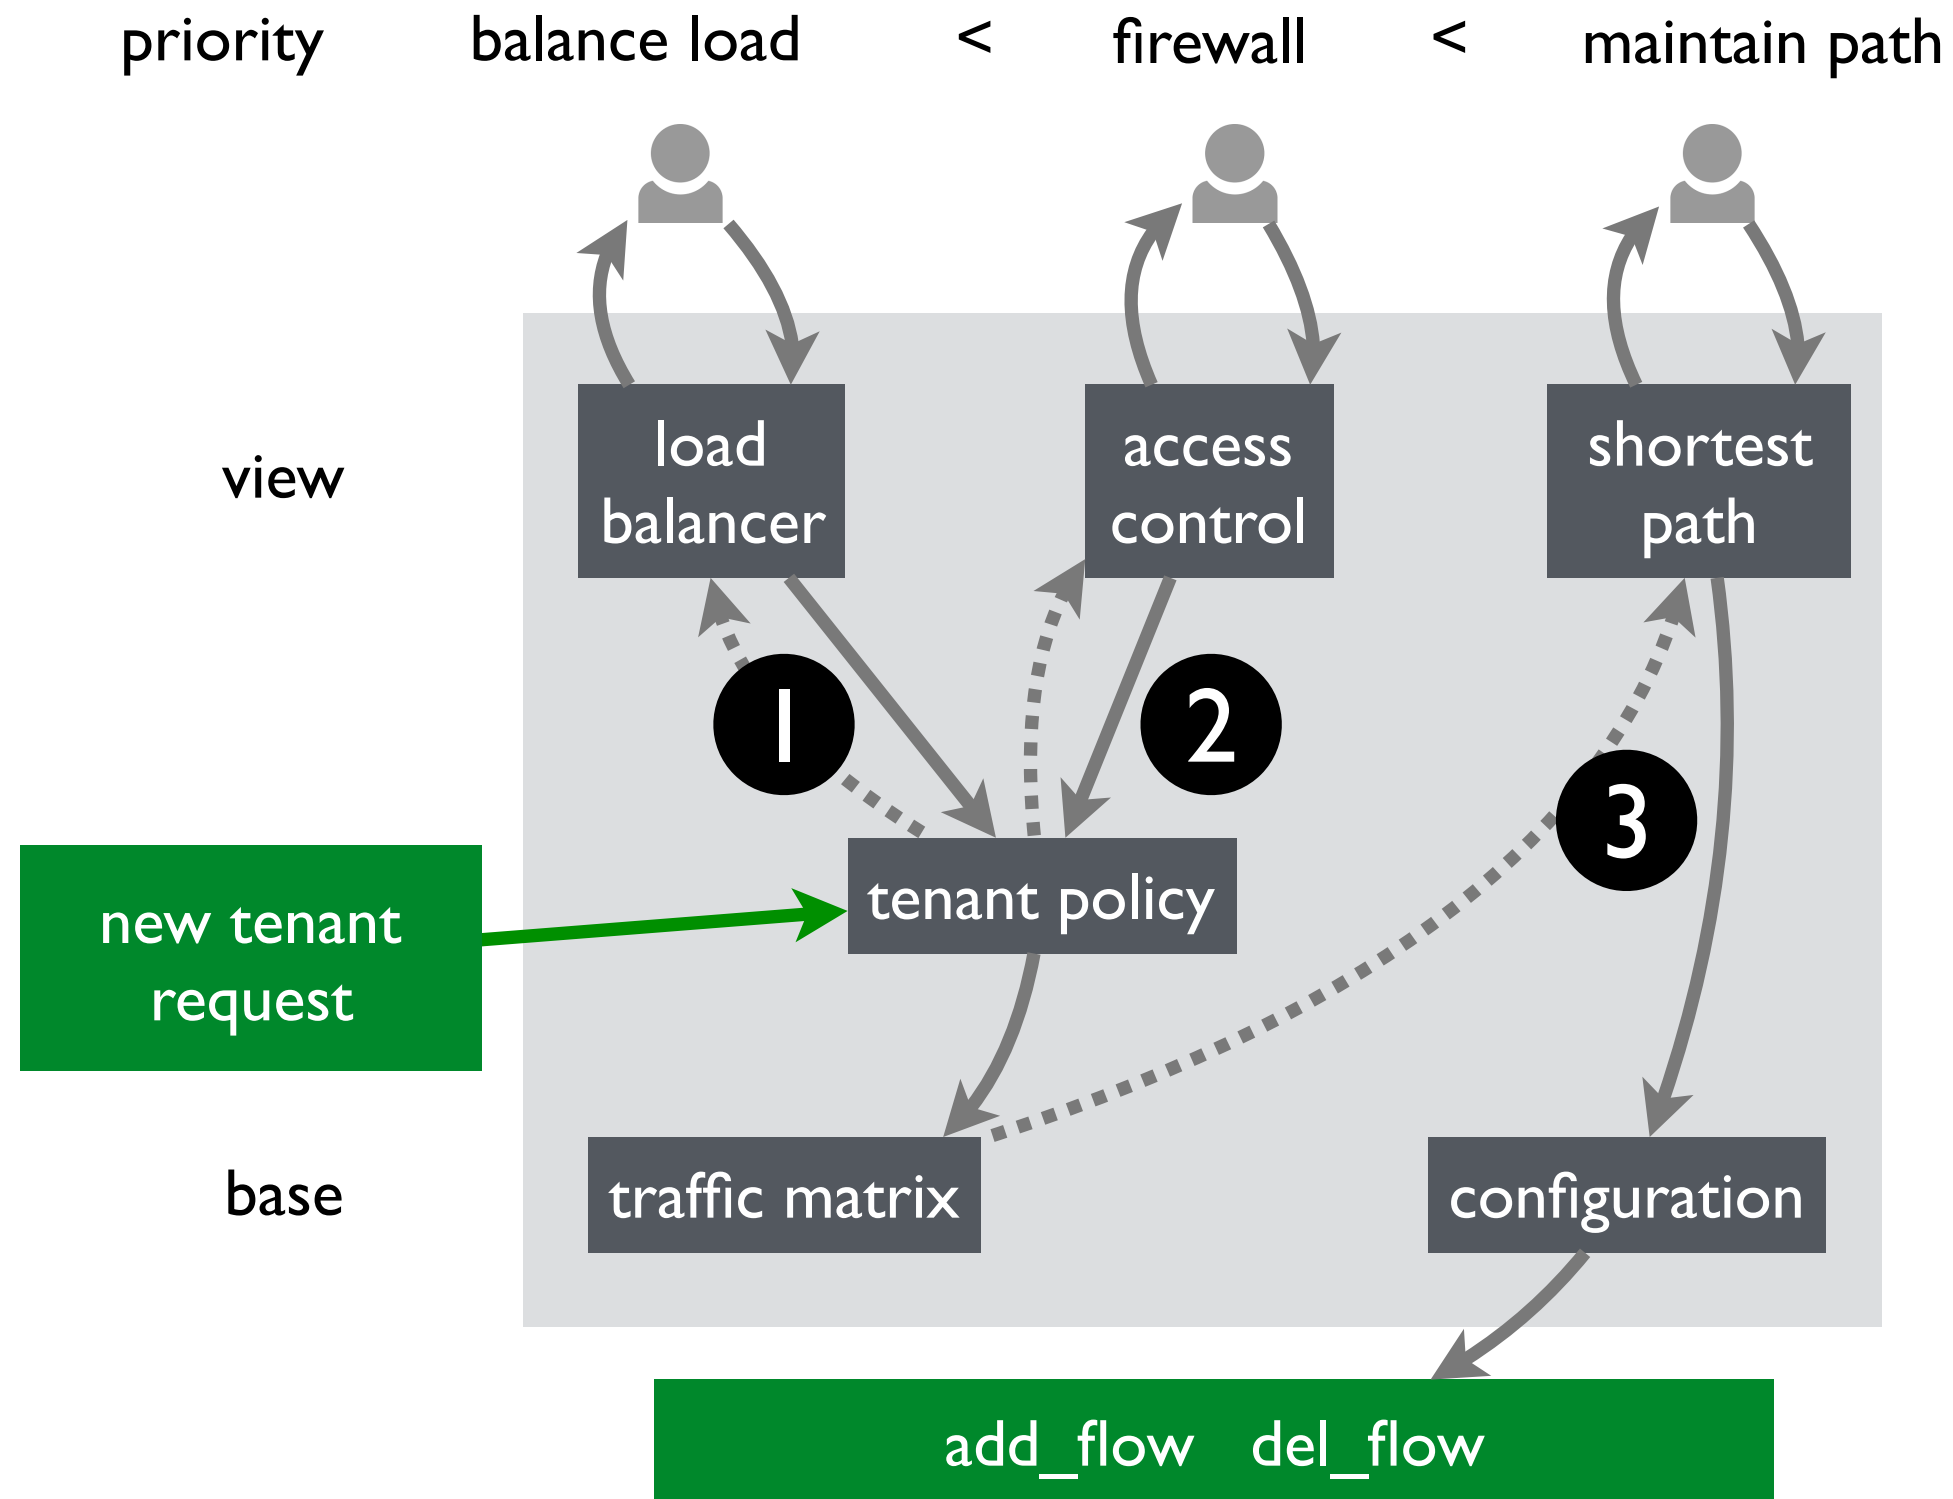
remove outdated inbound routes

```
CREATE RULE FW2 AS
  ON DELETE TO tm
  WHERE (SELECT count(*)
    FROM tm WHERE src = OLD.src AND
    dst = OLD.dst) = 1
  DO ALSO
    DELETE FROM FW
    WHERE end2 = OLD.src
    AND end1 = OLD.dst;
```

# orchestration across representations

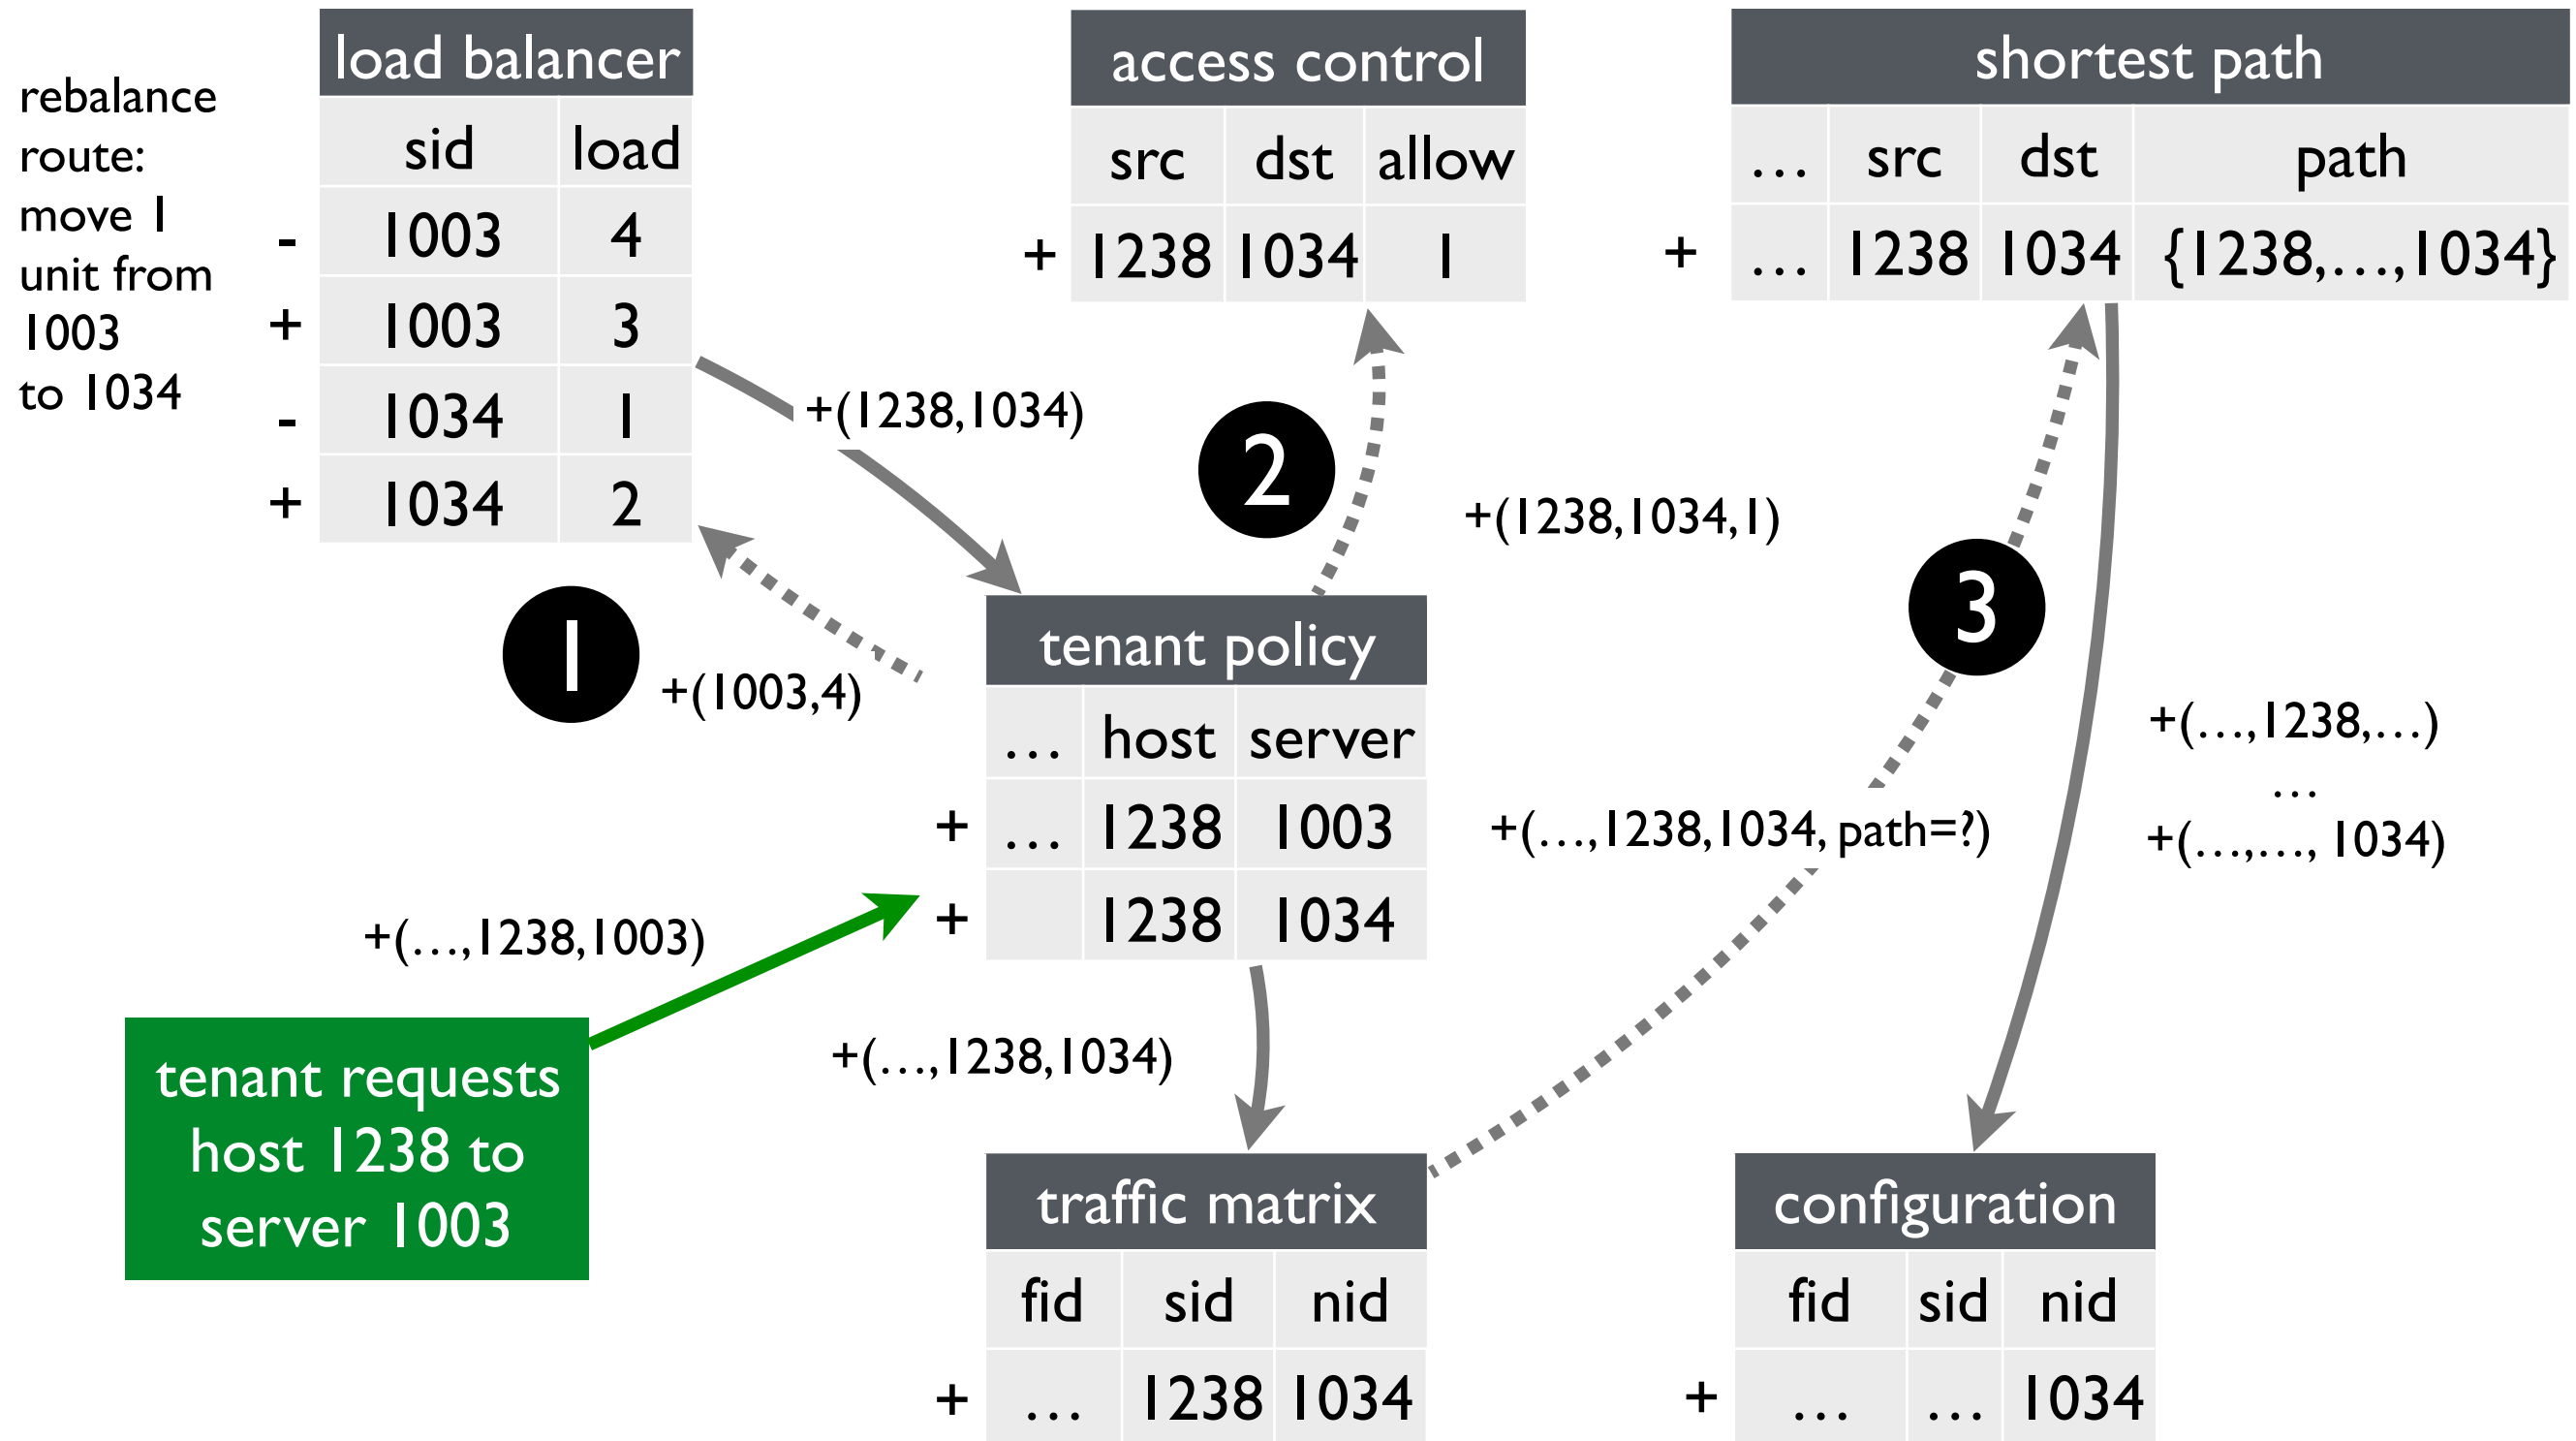


# orchestration across applications





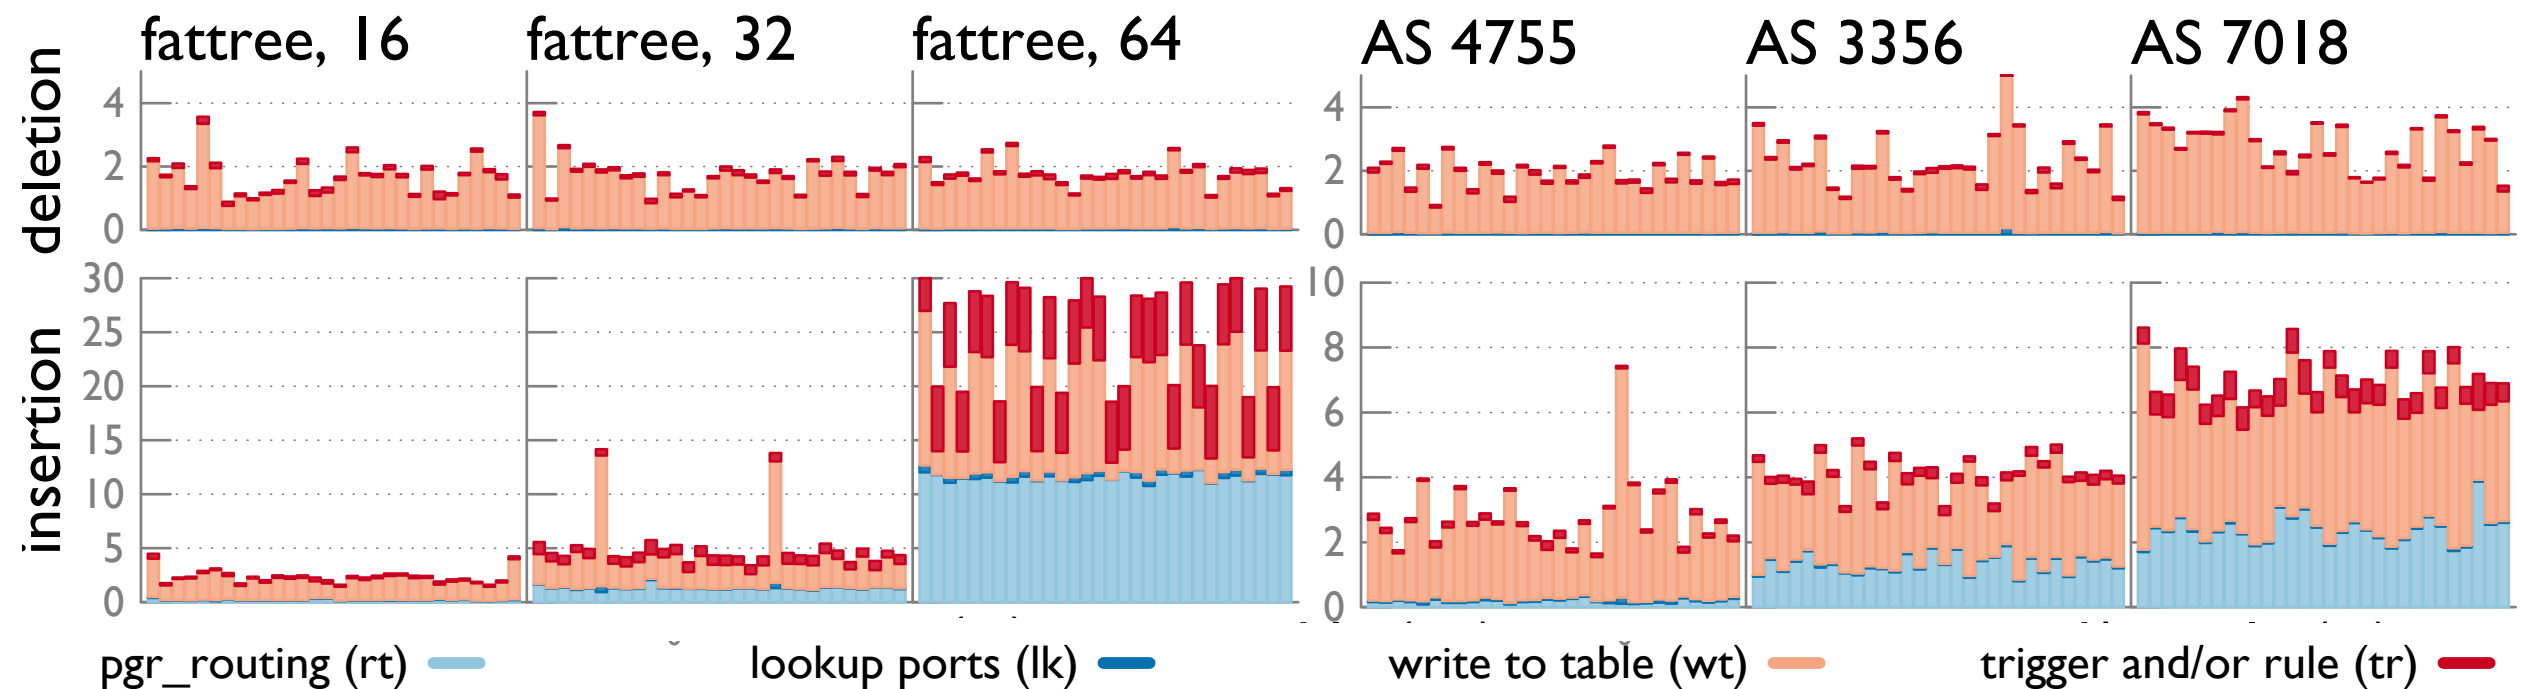
# orchestration across applications



orchestrated updates: install alternative route that is load-balanced and safe

# evaluation

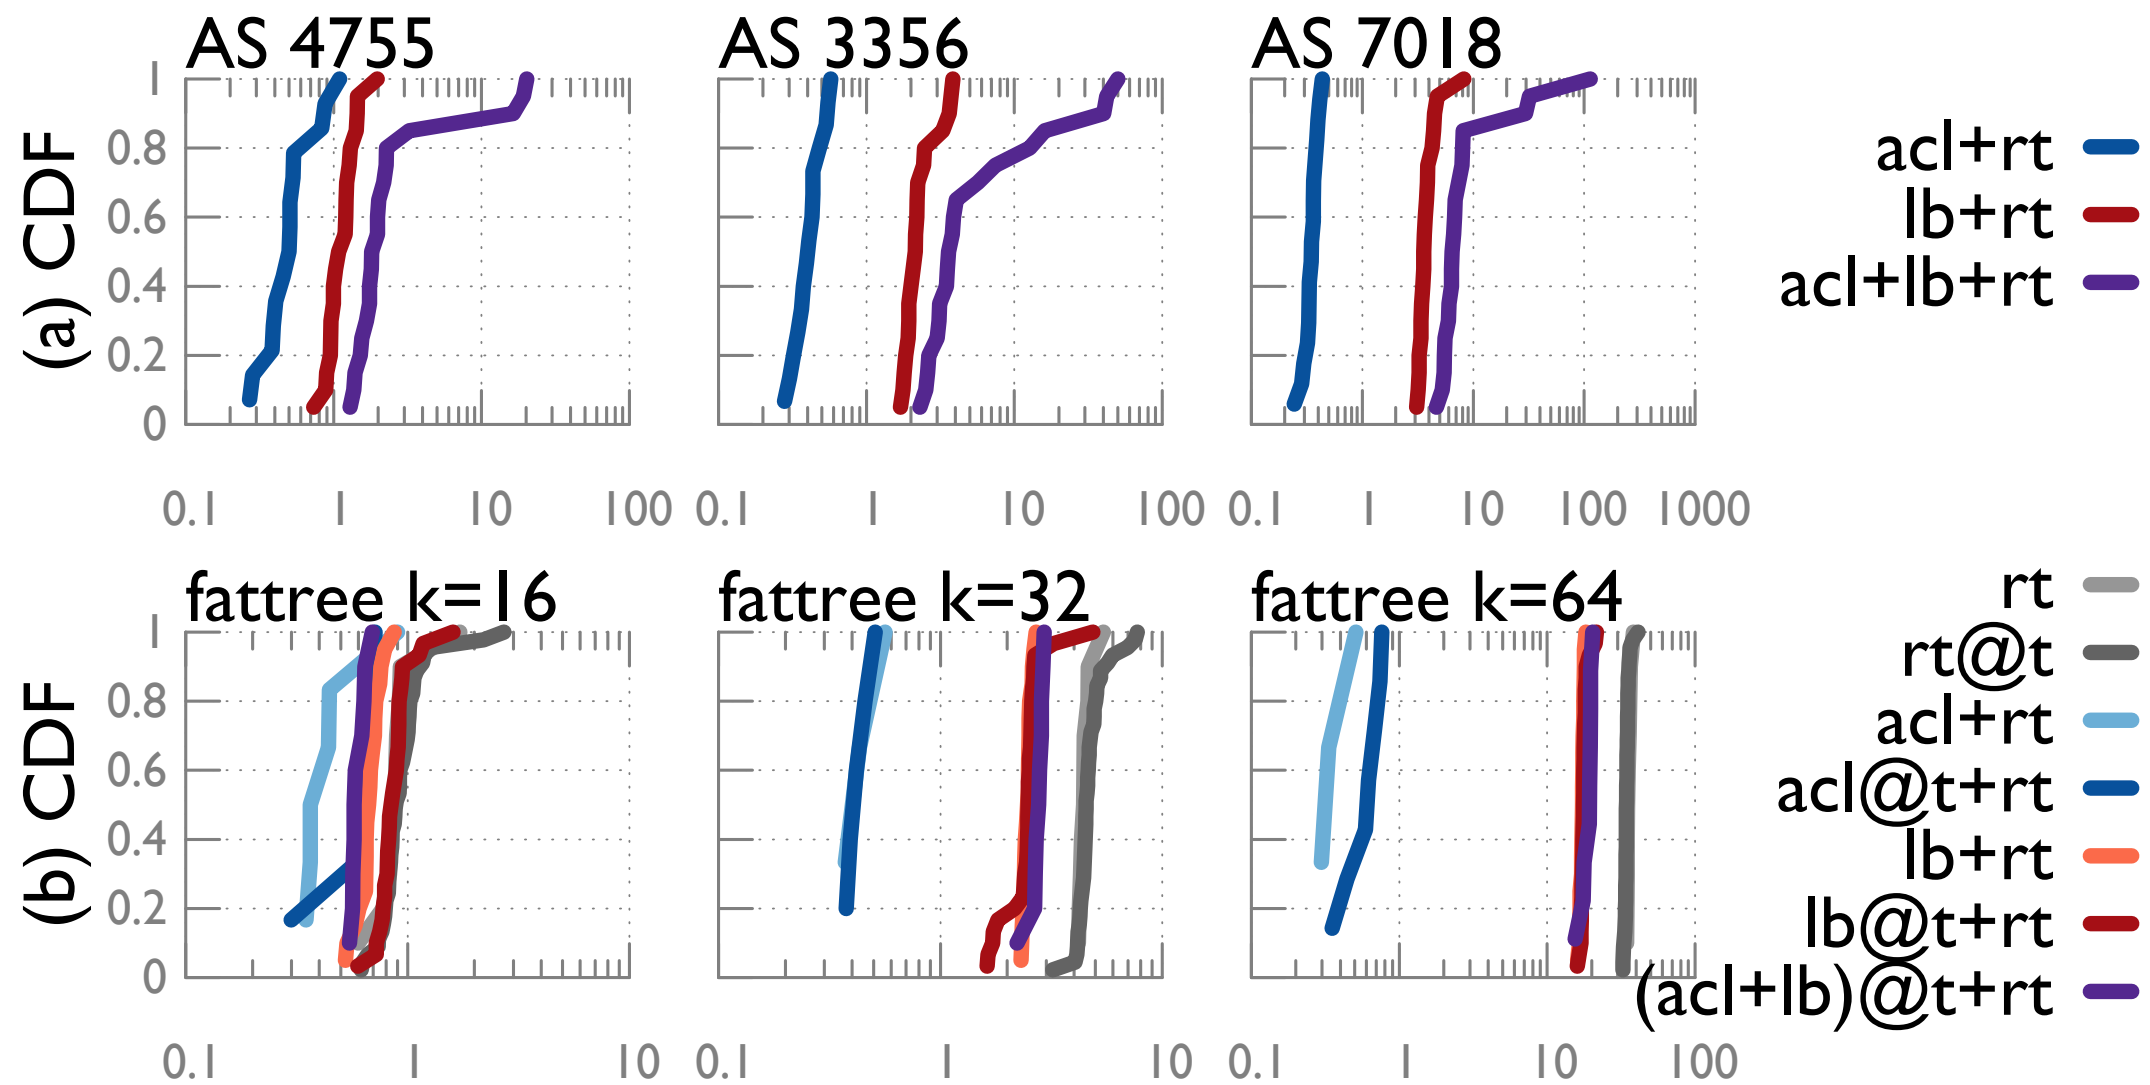
profiling database delay — route insertion/deletion



fat-tree			ISP		
k	switches	links	AS#	nodes	links
16	320	3072	4755	142	258
32	1280	24576	3356	1772	13640
64	5120	196608	7018	25382	11292
			2914	5939	16520

# evaluation

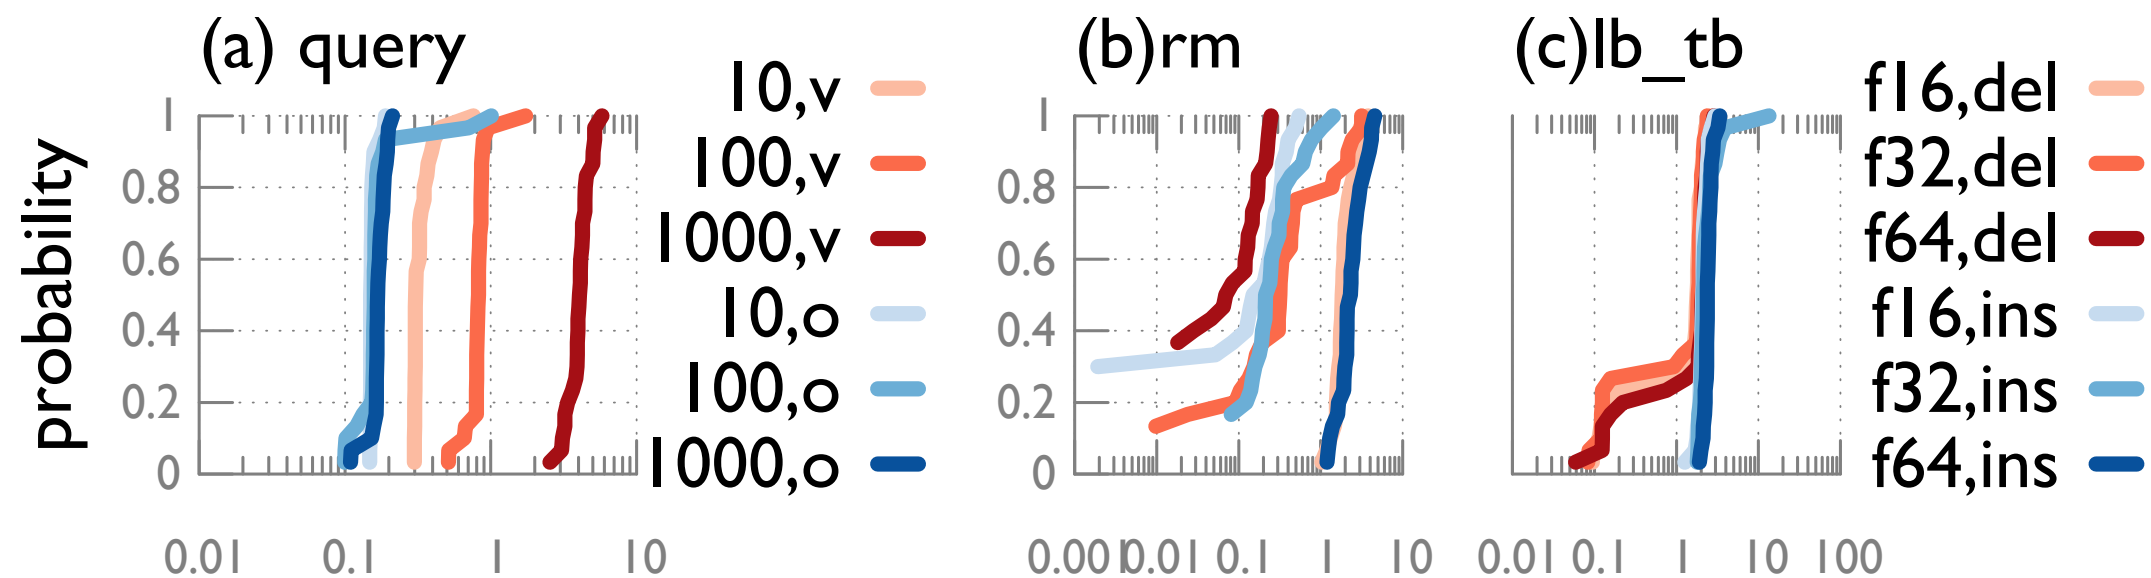
orchestrating access control(acl), load balancer(lb), and routing(rt): normalized per-rule delay (ms)



# evaluation

## optimizing application—materializing views

- faster access to materialized view (a)
- small maintenance delay (b,c)



# looking forward

use of standard SQL database enables direct application of many database theories and facilities

- revisit concurrency and recovery control
  - transaction processing
- revisit state distribution, interoperability (IXP)
  - distributed and federated database

## ongoing work

- bootstrapping database
- synthesizing orchestration

# thanks



# playtime

website (quick start, tutorials ...)

<https://ravel-net.org>

github

<https://github.com/ravel-net>

download *Ravel* (vm image)

<https://cis.temple.edu/~adw/ravel/ravelvm.zip>